


# Review of Computer Networks

---

Yao Liu

# How do applications talk to each other?



# Layered protocols

- Networks are organized as a series of layers, in order to reduce complexity
  - Each layer builds upon the one below it and offers services to the one above it.
  - Between each layer is an interface

# Physical and data link layers

- **Physical Layer:** Transmit and receive bits on physical media
  - analog and digital transmission
  - a definition of the 0 and 1 bits
  - bit rate (bandwidth)
- **Data Link Layer:** Provide error-free bit streams across physical media
  - Error detection/correction
  - reliability
  - flow control

# Network/Internet layer

Controls the operations of the network

Provides host-to-host connections

- **Internetworking** of both homogeneous and heterogeneous networks
- **Routing**: determining the path from the source of a message to its destination
- **Congestion Control**: handling traffic jams


# Transport layer

- Provides end-to-end connections
- **Packetization**: cut the messages into smaller chunks (packets)
- An ensuing issue is **ordering**: the receiving end must make sure that the user receives the packets in the right order
- End-to-end **flow control**

# Application layer

- file transfer, email, WWW, multimedia, and so on

# Internet Protocol Suite Reference Model





# Network / Internet layer

- Provides logical communication between hosts
- Internet Protocol (IP)
  - Key tool used today to build scalable, heterogeneous internetworks
  - It runs on all the nodes in a collection of networks
  - Hosts are assigned **IP addresses** based on the networks they locate in

# IPv4 addresses

- Fixed-length, 32-bit address
- Each network has its network ID
- Every interface in that network has an IP address comprising the network ID and a host ID
- *IP addresses identify interfaces*
  - There is no 1-to-1 correspondence between IP addresses and nodes (hosts or routers)


# IPv4 address notation

- A **dotted decimal** notation is used by human
- Binary: 11000000 01111111 11111101 00000001  
Decimal: 192. 127. 253. 1
- Examples:
  - cs.binghamton.edu = 128.226.118.12
  - remote02.cs.binghamton.edu = 128.226.114.202

# Packet forwarding

- Each host has a default router
- Each router maintains a **forwarding table**
- Every packet contains destination's address
- If directly connected to destination network, then forward to host
- If not directly connected to destination network, then forward to some router
- Forwarding table maps **network number** into **next hop**

# Packet forwarding



NetworkNum	NextHop
1	R1
2	Interface 1
3	Interface 0
4	R3

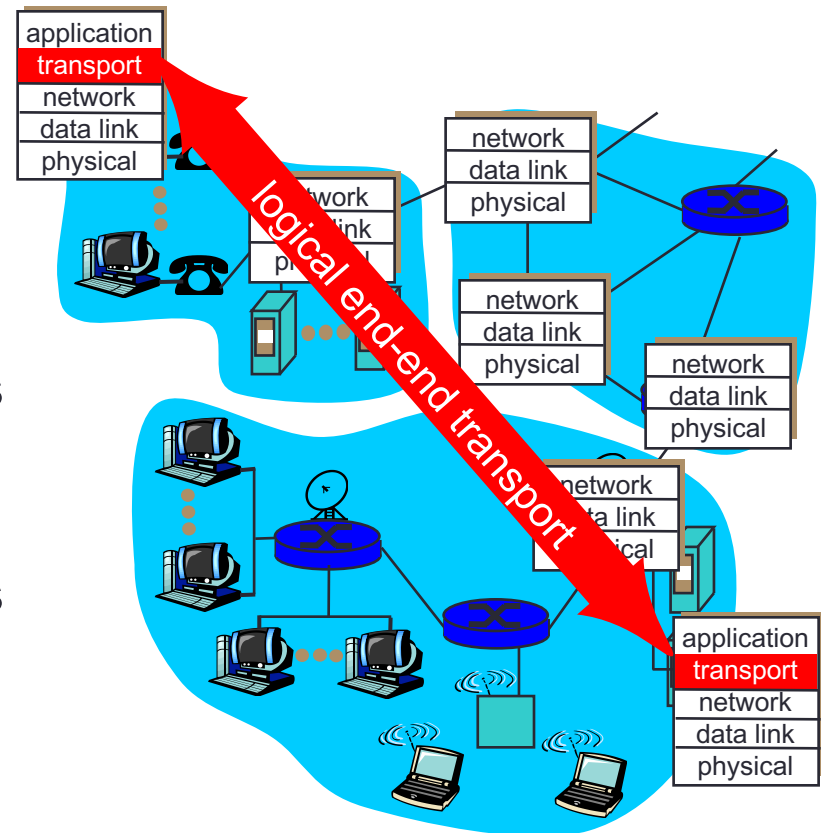
R2's forwarding table. Two are direct routes

# IPv4 packet header

0	4	8	12	16	19	24	31
Vers		HLen		Service Type		Total Length	
Identification				Flags		Fragment Offset	
Time to Live		Protocol		Header Checksum			
Source IP Address							
Destination IP Address							
IP Options						Padding	
Data							
• • • •							

# Transport services and protocols

- Provide *logical communication* between **application processes** running on different hosts
- Transport protocols run in end systems
  - sender side: breaks application messages into **segments**, passes to network layer
  - receiver side: reassembles segments into messages, passes to application layer
- More than one transport protocol available to applications
  - Internet: TCP and UDP



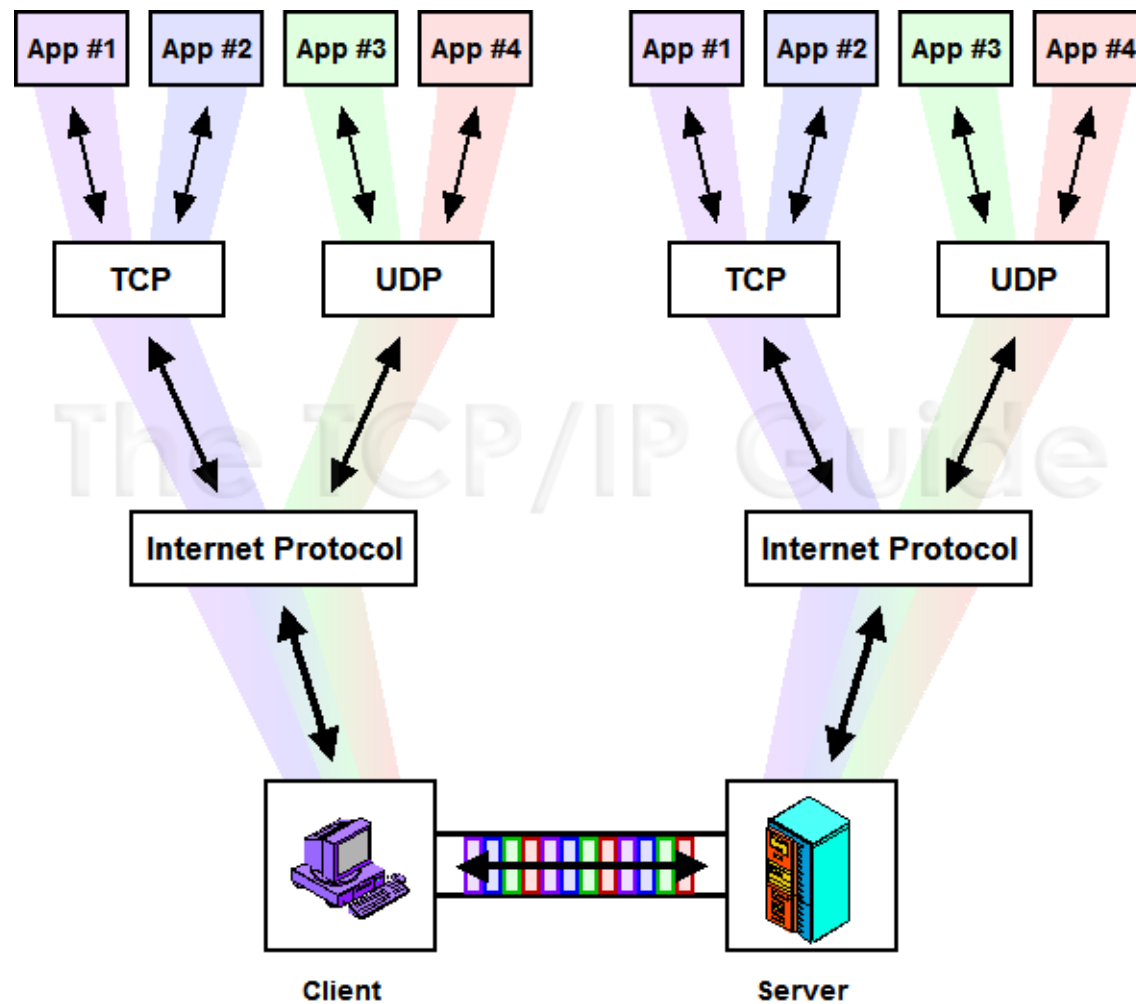
# Transport layer

- Basic service: moving data from one application (or process) to a remote application
- Functions:
  - Packetization and reassembly of application data
  - Connection management
  - Error control
  - Reliability
  - Flow control
  - Congestion control (on the Internet)



# Transport-layer protocols

- **Reliable, in-order** delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- **Unreliable, unordered** delivery (UDP)
  - none of above
- Services not available:
  - delay guarantees
  - bandwidth guarantees



# Socket

- The end-point where a local application process attaches to the network is called a socket
- End point determined by two things:
  - Host address: IP address is *Network Layer*
  - Port number: is *Transport Layer*
- Two end-points determine a connection: socket pair
  - e.g., 128.226.180.163,port 21 + 198.69.10.2,port 1500
  - e.g., 128.226.180.163,port 21 + 198.69.10.2,port 1499

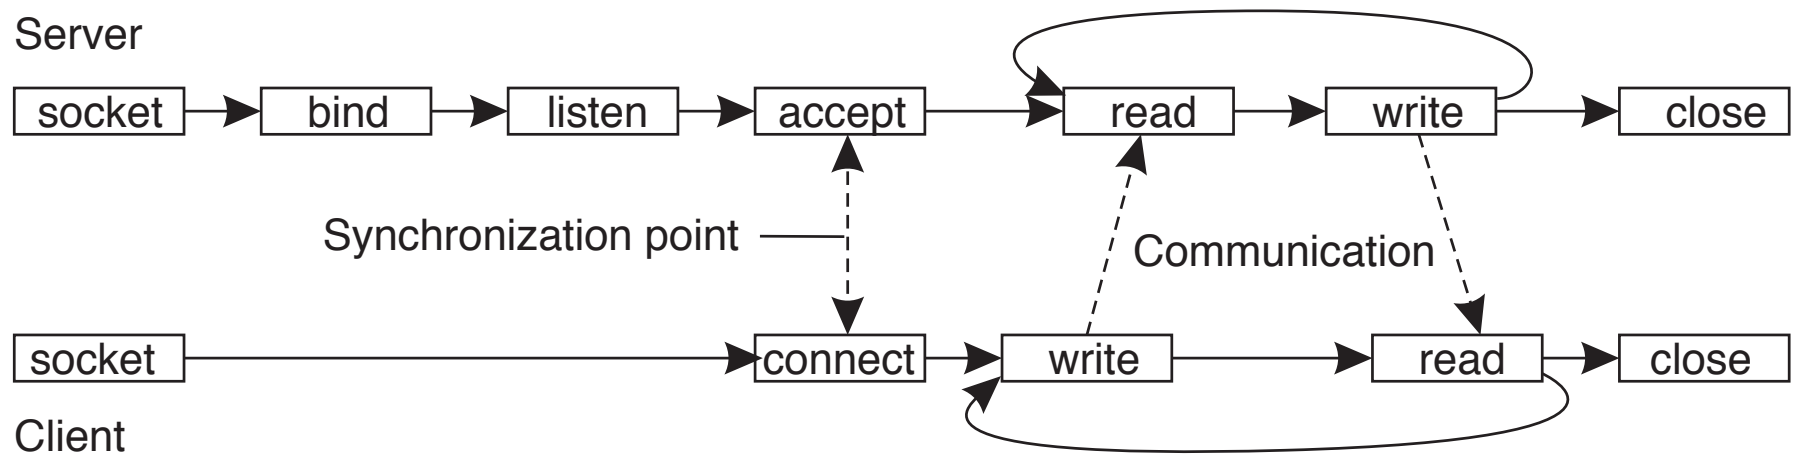
# Port numbers

- Basic mechanism for multiplexing applications per host
  - 65,535 ports available
- Ports <1024 are reserved
  - Only privileged processes (e.g., superuser) may access
  - Well-known , reserved services (see /etc/services in Unix):
    - ftp 21/tcp, ssh 22/tcp, telnet 23/tcp, http 80/tcp, ntp 123/tcp , snmp 161/udp
- “I tried to open a port and got an error”
  - Port collision: only one application per port per host
  - Dangling sockets...

# Socket API

- Socket is the interface between an application and the network
- The interface defines operations for:
  - Creating a socket
  - Attaching a socket to the network
  - Sending and receiving messages through the socket
  - Closing the socket

# TCP socket API



# Socket API

```
int sockfd = socket(address_family, type, protocol);
```

- The socket number returned is the socket descriptor for the newly created socket

```
• int sockfd = socket (PF_INET, SOCK_STREAM, 0);
```

```
• int sockfd = socket (PF_INET, SOCK_DGRAM, 0);
```

The combination of PF\_INET and SOCK\_STREAM implies TCP

The combination of PF\_INET and SOCK\_DGRAM implies UDP

# Socket API

- Passive Open (on server)

int **bind**(int socket, struct sockaddr \*addr, int addr\_len)

int **listen**(int socket, int backlog)

int **accept**(int socket, struct sockaddr \*addr, int addr\_len)

- Active Open (on client)

int **connect**(int socket, struct sockaddr \*addr, int addr\_len)

- Sending/Receiving Messages

int **send**(int socket, char \*msg, int mlen, int flags)

int **recv**(int socket, char \*buf, int blen, int flags)



- The socket API offered by the Linux operating system is in C.
- But many programming languages also provide similar interfaces, providing accesses to the socket interface offered by the operating system.

# Multiplexing/Demultiplexing

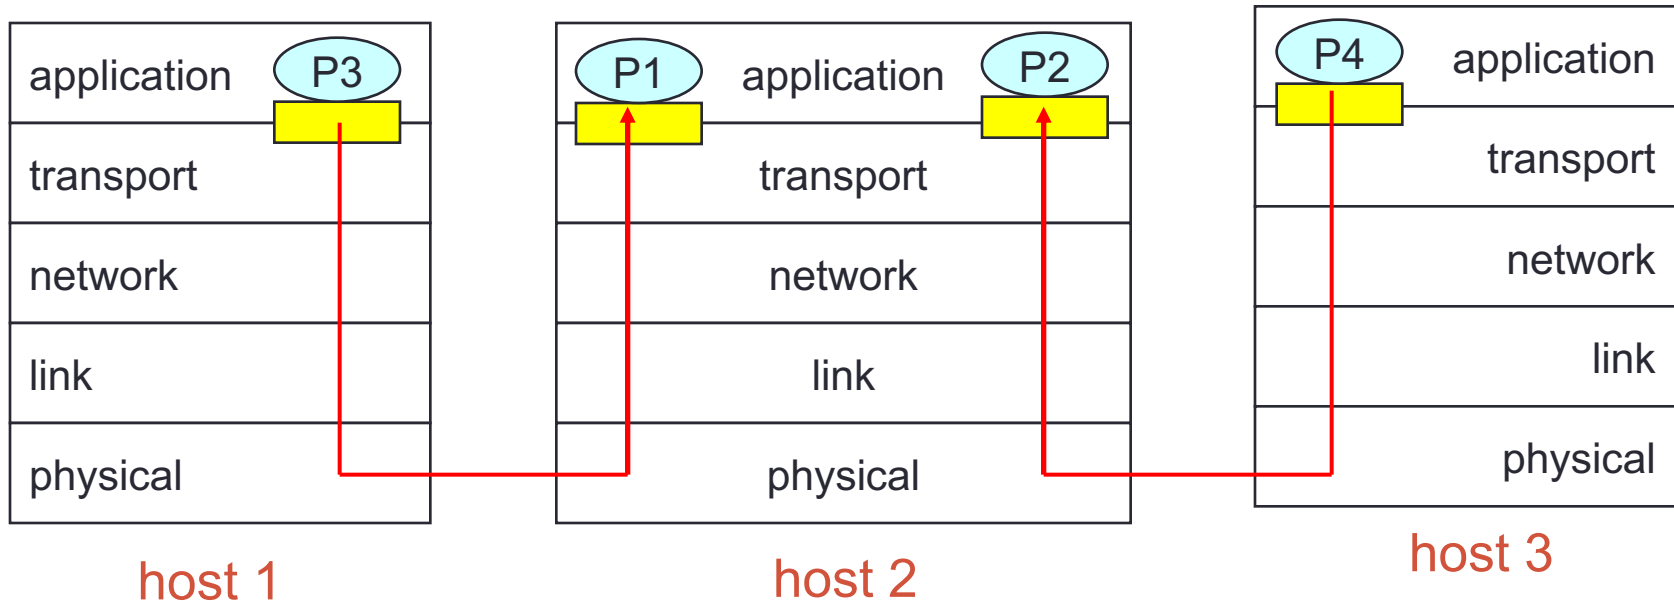
## Demultiplexing at recv host:

delivering received segments  
to correct socket

## Multiplexing at send host:

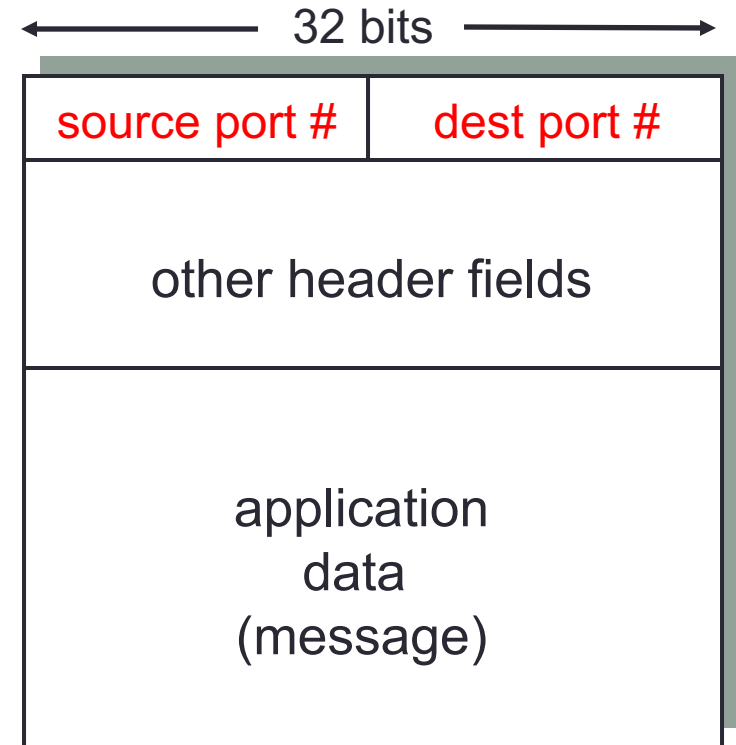
gathering data from multiple  
sockets, enveloping data with  
header (later used for  
demultiplexing)

 = socket       = process



# How demultiplexing works

- Host receives IP packets
  - each packet has source IP address, destination IP address
  - each packet carries one transport-layer segment
  - each segment has source, destination **port numbers**
- Host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format


# UDP

- “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- **connectionless**:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

## Why is there UDP?

- **no connection establishment** (which can add delay)
- **simple**: no connection state at sender, receiver
- **small** segment **header**
- **no congestion control**: UDP can blast away as fast as desired


# UDP format



# Connectionless demultiplexing

- **UDP** socket identified by
  - (dest IP address, dest port number)
- When host receives UDP segment:
  - checks destination port number in segment
  - directs UDP segment to socket with that port number
- IP datagrams with *different* source IP addresses and/or source port numbers directed to *same* socket

# Connectionless demultiplexing



SP provides “return address”

# TCP


- TCP provides
  - reliable, stream delivery service
  - full-duplex (bidirectional) communication
  - flow control
  - congestion control
  - connection establishment and destruction



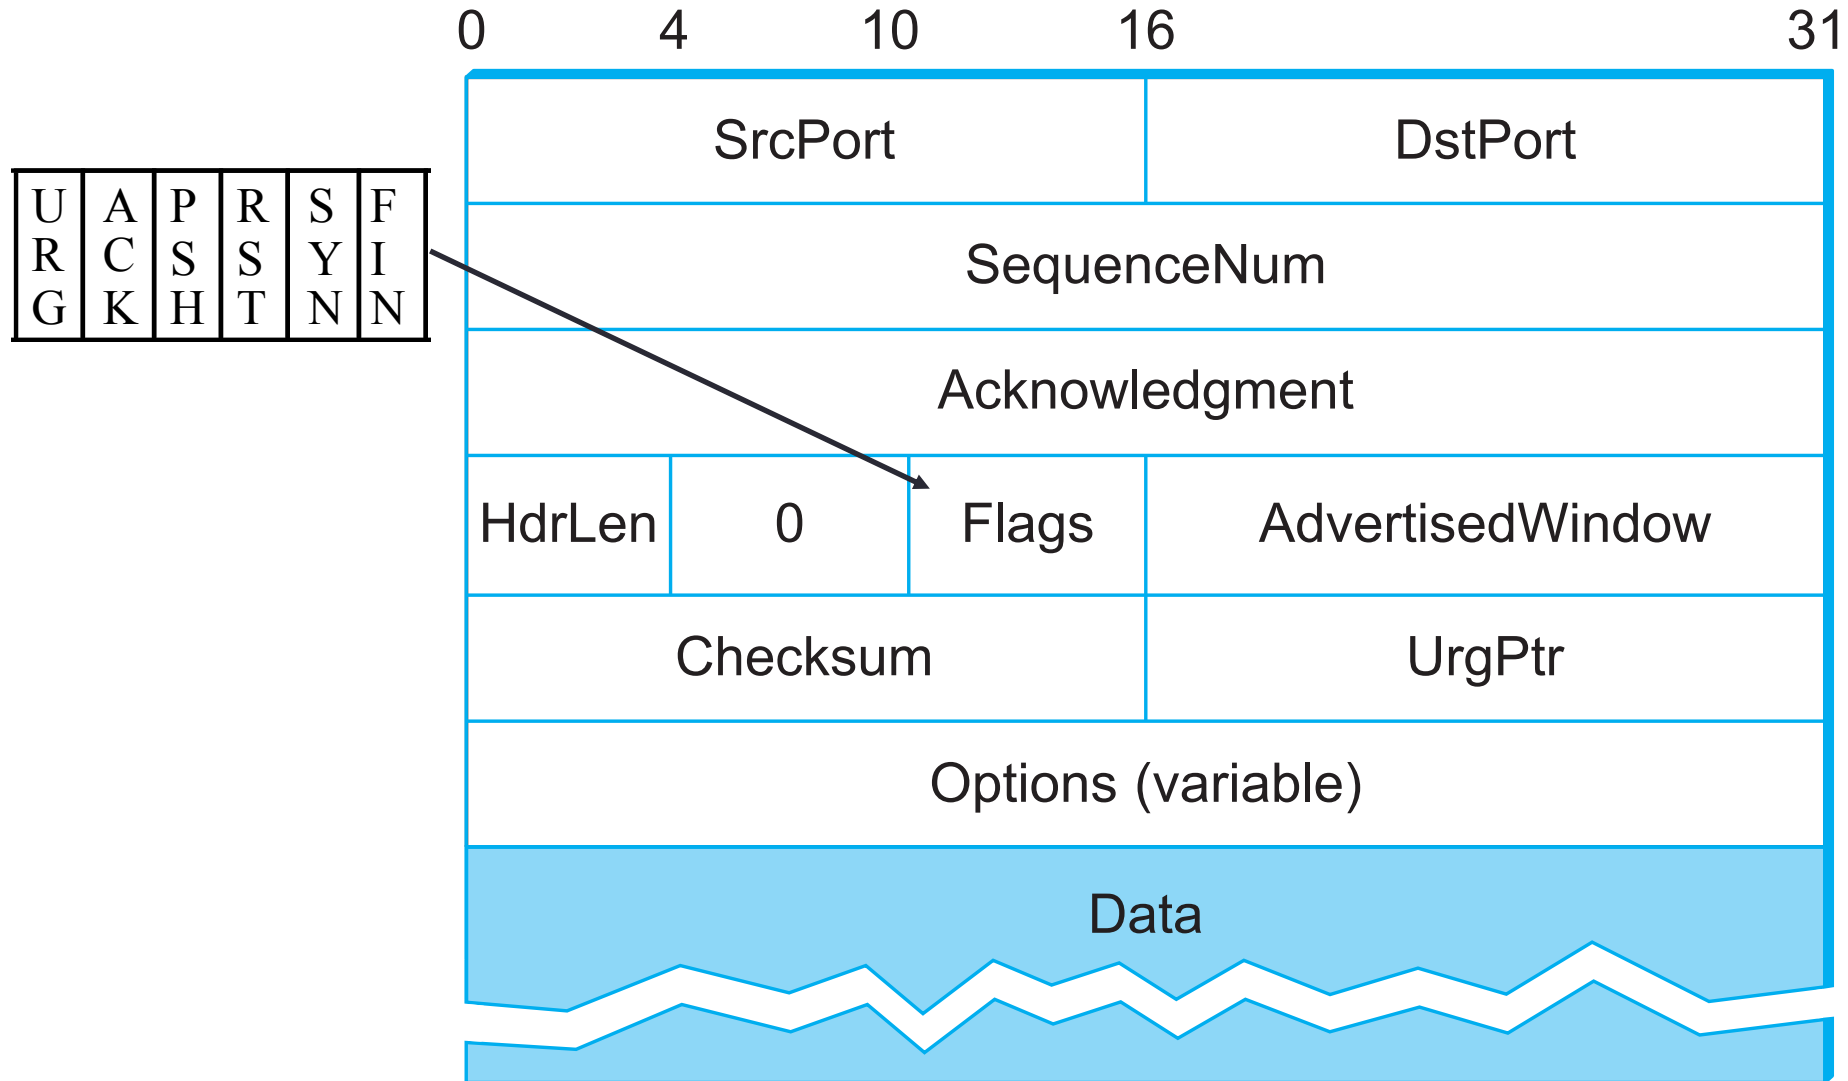
# *Connection-oriented* demultiplexing

- **TCP** socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- Receiving host uses all four values to direct segment to appropriate socket
- Server hosts may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client

# Connection-oriented demultiplexing



# TCP segment format



# Stream delivery service

- Data from an application is treated as a stream of bytes
- The stream is divided into **segments** for delivery.
  - this division is up to TCP; application data boundaries are ignored
- Conceptually, **every byte** has a sequence #.
- Only the sequence # of the first byte of the segment is sent.

# Three-way hand shaking

- Process A initiates a connection to process B.
- A sends a segment with ISN  $x$ , chosen according to A's clock, and with SYN bit set to 1, this is the **connection request** message.

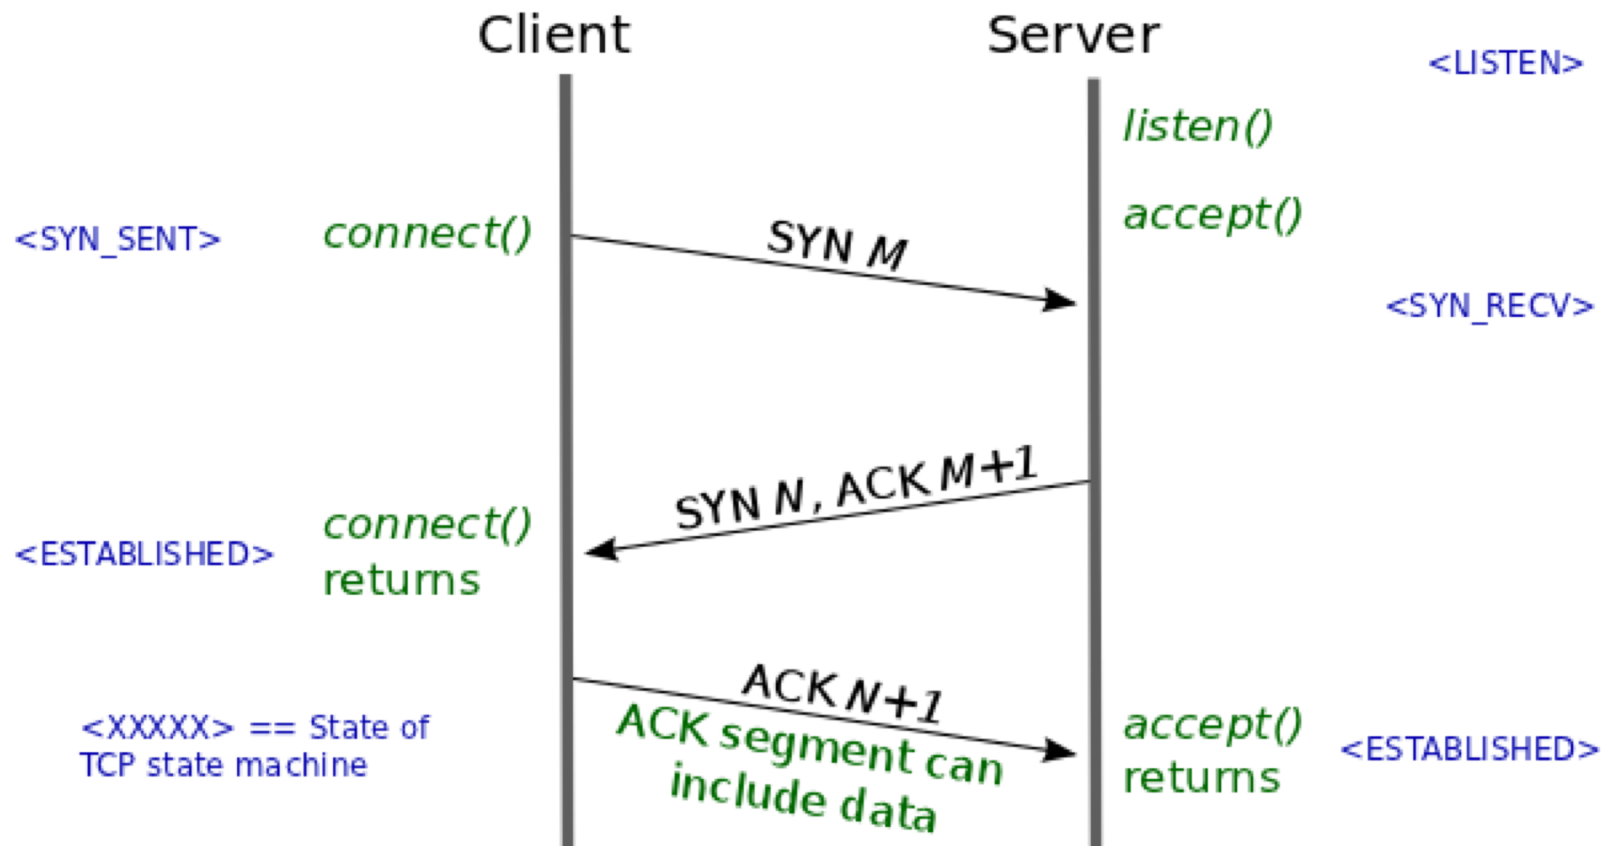
# Three-way hand shaking

- B returns a segment whose  $\text{ack} = x + 1$ ,  $\text{seq} = y$ , the ISN of B chosen according to B's clock, and  $\text{SYN} = 1$ .
  - this is the **connection acceptance** message.
- After receiving the acceptance, A considers the connection established.

# Three-way hand shaking

- A then sends B a segment with  $\text{ack}=y+1$  and  $\text{SYN}=0$ .
  - this is the **connection confirmation** message
  - all segments from this point on will have  $\text{SYN}=0$
- After receiving the confirmation, B considers the connection established.

# Visualize It

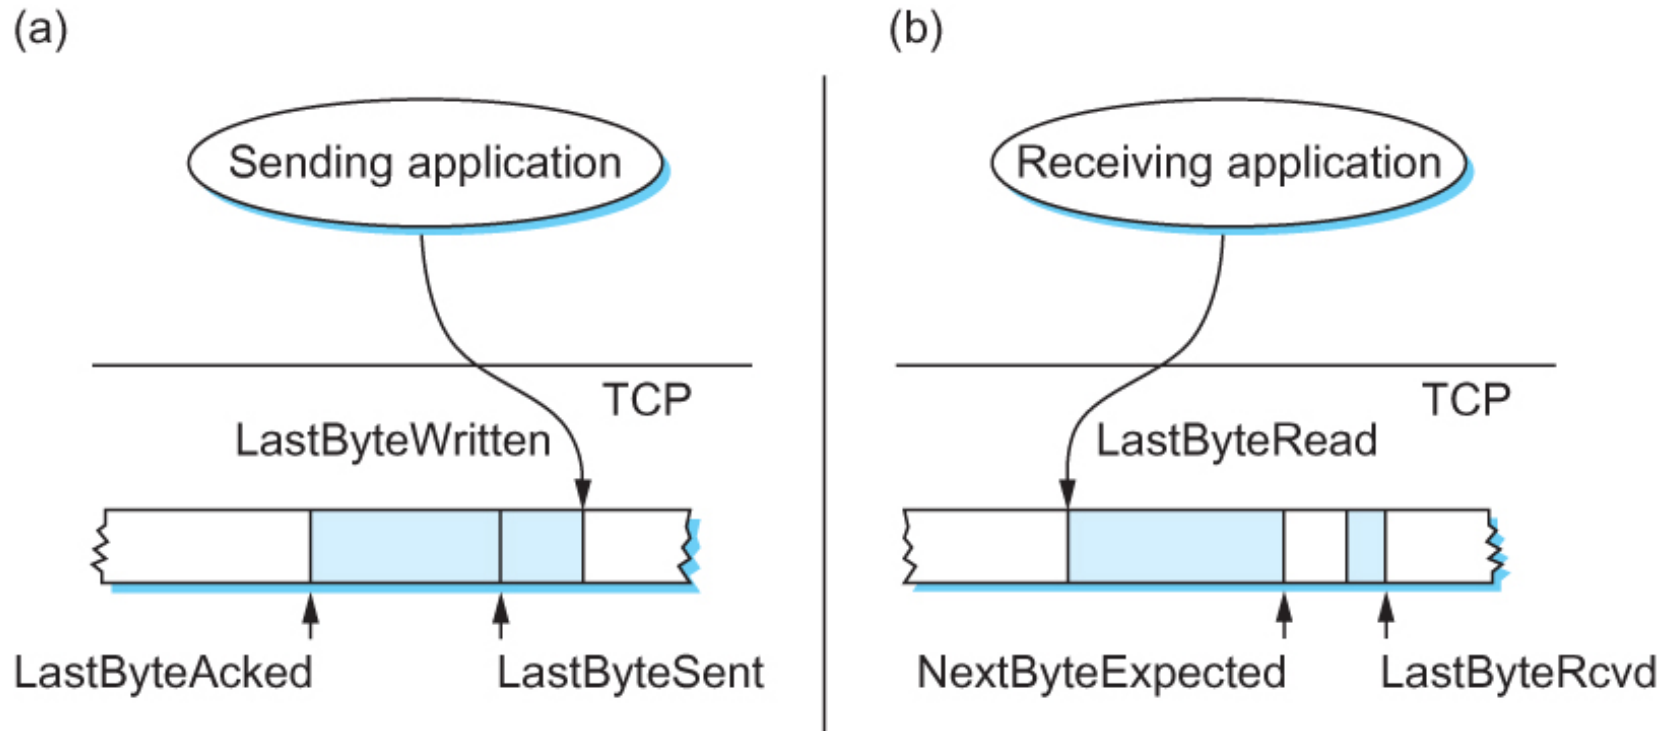




# TCP acknowledgments

- TCP acknowledges the next byte expected (not the last one received).
- $\text{Ack}=x$  indicates *all* bytes up to and including  $x-1$  has been successfully received.
- Also, a pure Ack is simply a segment with no data.
- What if a segment is received out of order? (next slide)

# Sliding Window in TCP



Relationship between the TCP send buffer (a) and the receive buffer (b).

# TCP flow control

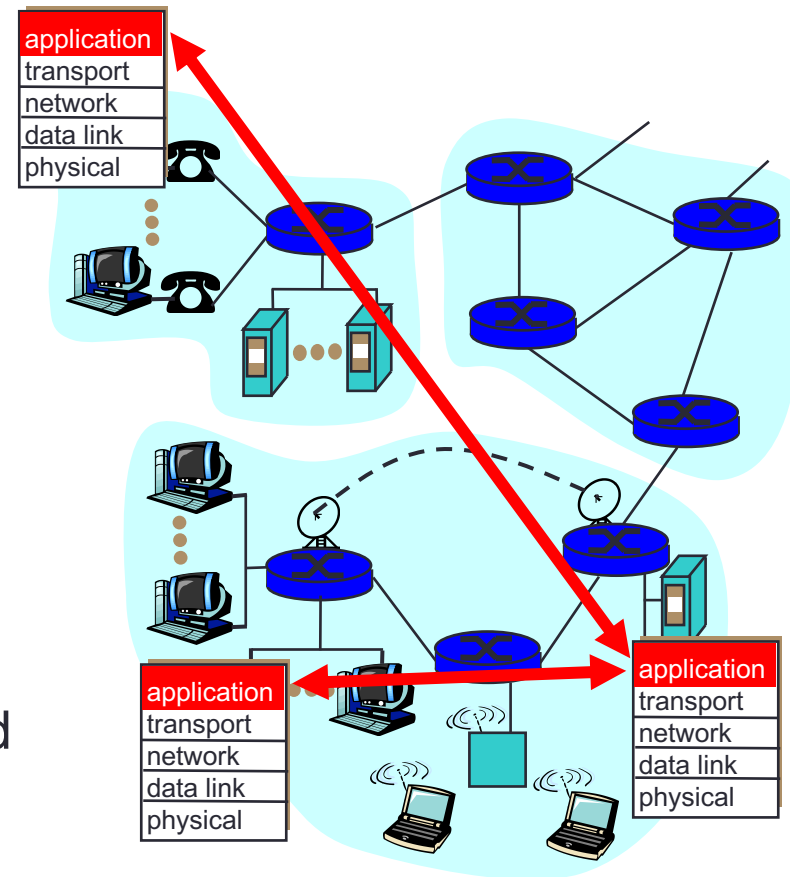
- Flow control aims to eliminate the possibility of overflowing the receiver's buffer.
- Receiver side:
  - $\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$
  - **AdvertisedWindow** =  $\text{MaxRcvBuffer} - ((\text{NextByteExpected} - 1) - \text{LastByteRead})$
- Sender side:
  - **$\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertisedWindow}$**

# TCP congestion control

- When load builds up in a network, congestion builds up.
- Senders limit the rate at which it sends traffic as a function of the perceived network congestion.
- If there is little to no perceived congestion,
  - Increase the sending rate
- If congestion is detected,
  - Decrease the sending rate
- Detect congestion via:
  - Lost packets

# Applications and application-layer protocols

- Application: communicating, distributed processes
  - e.g., e-mail, Web, instant messaging
  - running in end systems (hosts)
  - exchange messages
- Application-layer protocols
  - one “piece” of an application
  - define messages exchanged by applications and actions taken
  - use communication services provided by lower layer protocols (TCP, UDP)



# Application-layer protocol defines

- Types of messages exchanged, e.g., request & response messages
- Syntax of message types: what fields in messages & how fields are delineated
- Semantics of the fields, i.e., meaning of information in fields
- Rules for when and how processes send & respond to messages

## Public-domain protocols:

- Defined in RFCs, allows for interoperability
- e.g., HTTP, FTP

## Proprietary protocols:

- e.g., Skype

# Client-server paradigm

- Typical network app has two pieces: **client** and **server**
- **Client:**
  - initiates contact with server (“speaks first”)
  - typically requests service from server
  - e.g., web client implemented in browser
- **Server:**
  - provides requested service to client
  - e.g., web server sends requested web page

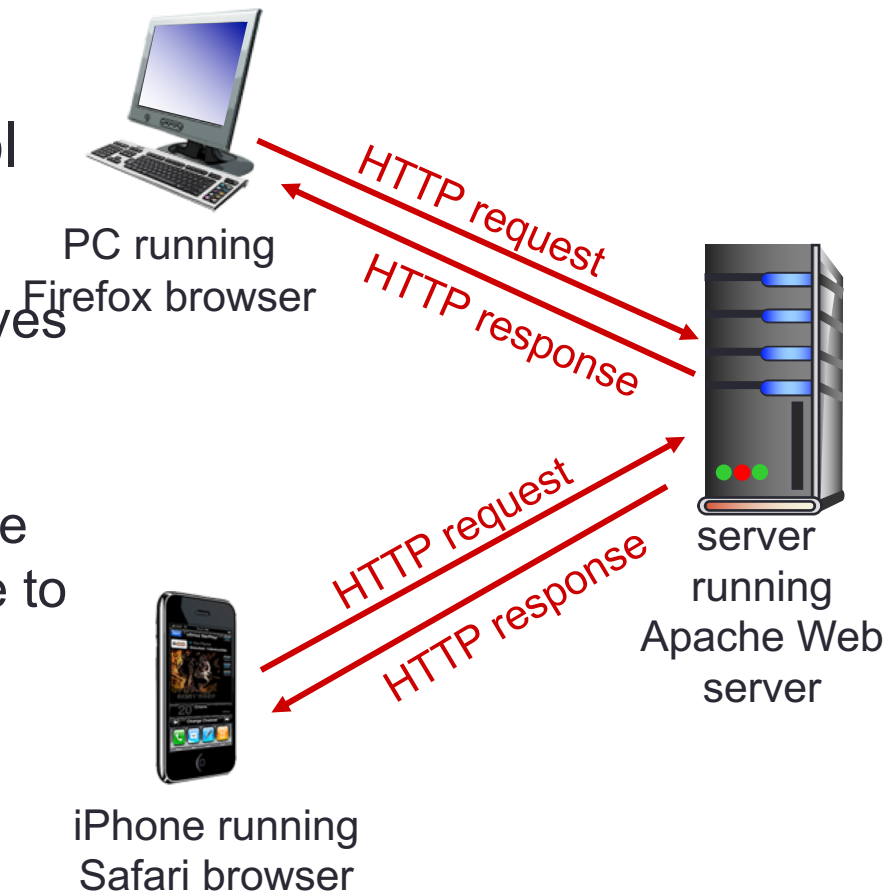
# World Wide Web

- Server provides access to web pages.
- Browsers request web pages from server using Hypertext Transfer Protocol (HTTP).
- Objects on the web are identified by Uniform Resource Locators (URL).



# HTTP

- A simple client-server protocol
- Web's application layer protocol
- Client/server model
  - **client**: browser that requests, receives (using the HTTP protocol), and displays Web objects
  - **server**: Web server sends (using the HTTP protocol) objects in response to requests
- HTTP/1.0: RFC 1945
- HTTP/1.1: RFC 2616
- HTTP/2: RFC 7540



# HTTP

- **Uses TCP:**
  - client initiates TCP connection to server, typically port 80
  - server accepts TCP connection from client
  - HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
  - TCP connection closed
- **HTTP is “stateless”**
  - server maintains no information about past client requests
- **Protocols that maintain “state” are complex!**
  - past history (state) must be maintained
  - if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# Web and HTTP

- **Web page** consists of **objects**
- Object can be HTML file, JPEG image, Javascript, audio file, video file, ...
- Web page consists of **base HTML-file** which includes several referenced objects
- Each object is addressable by a **URL**
- Example URL:

`www.someschool.edu/someDept/pic.gif`

host name

path name

# HTTP connections

- **Non-persistent HTTP**
  - At most one object is sent over a TCP connection.
  - HTTP/1.0 uses non-persistent HTTP
- **Persistent HTTP**
  - Multiple objects can be sent over single TCP connection between client and server.
  - HTTP/1.1 uses persistent connections in default mode

# Non-persistent HTTP

Suppose user enters URL (contains text, references to 10 jpeg images)

`www.someSchool.edu/someDepartment/index.html`

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

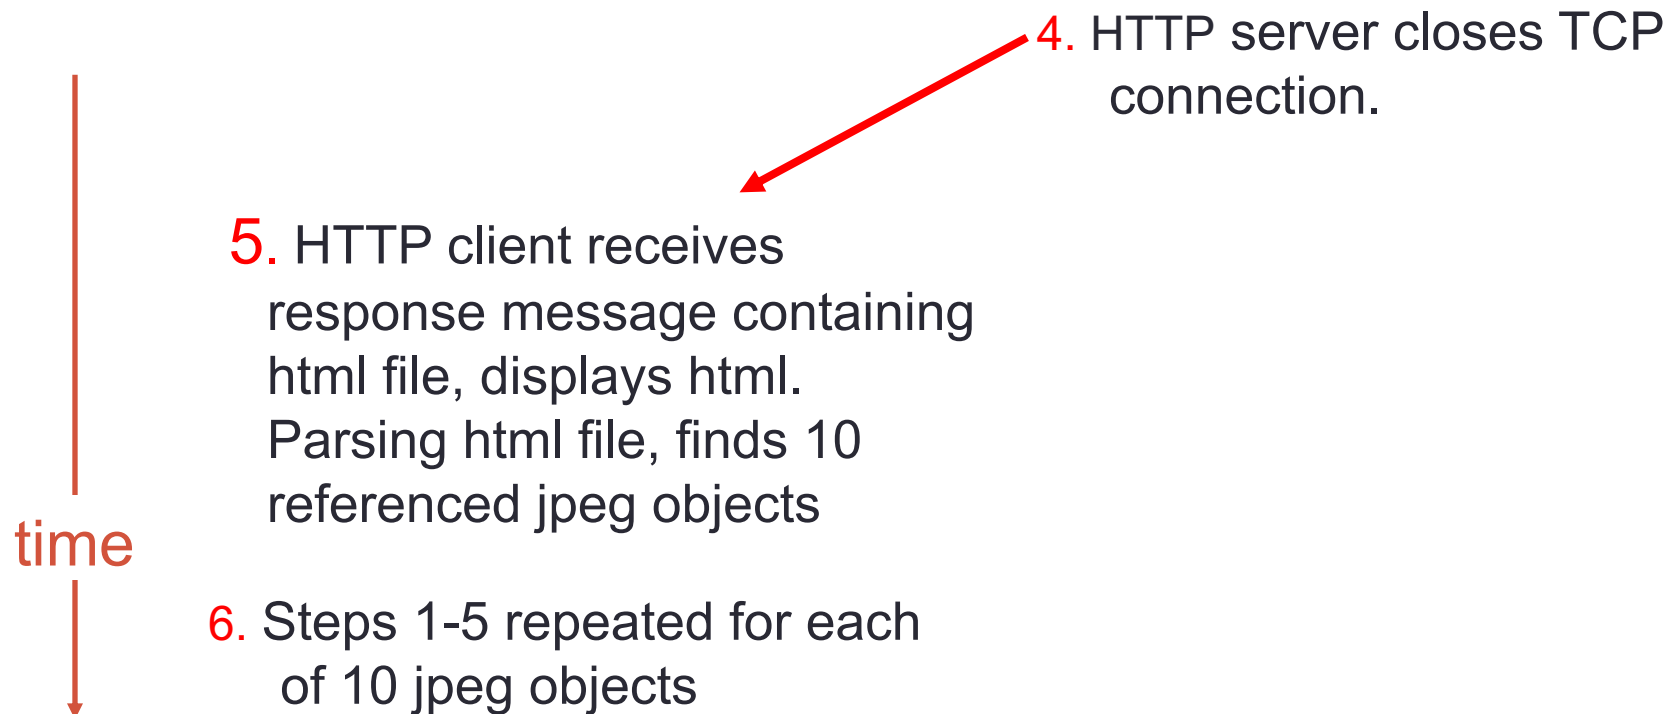
1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `/someDepartment/index.html`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

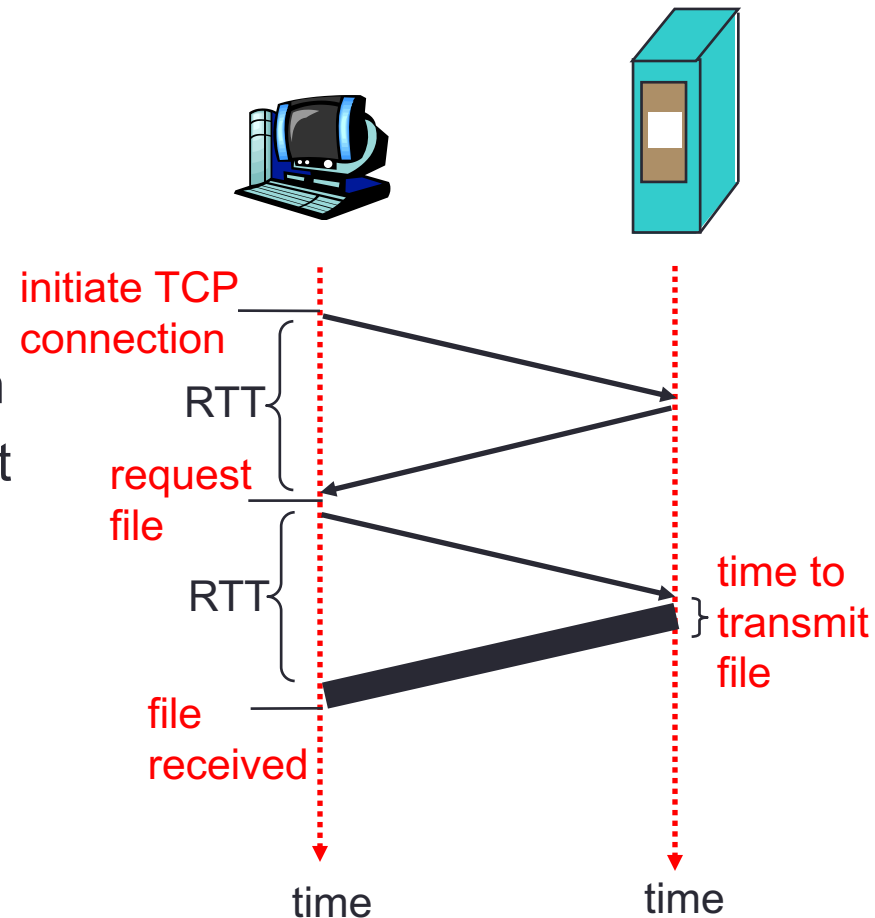
time





# Modeling response time

- **Definition of RTT:** time to send a small packet to travel from client to server and back.
- **Response time:**
  - one RTT to initiate TCP connection
  - one RTT for HTTP request and first few bytes of HTTP response to return
  - file transmission time
- **total =  $2RTT + \text{transmit time}$**


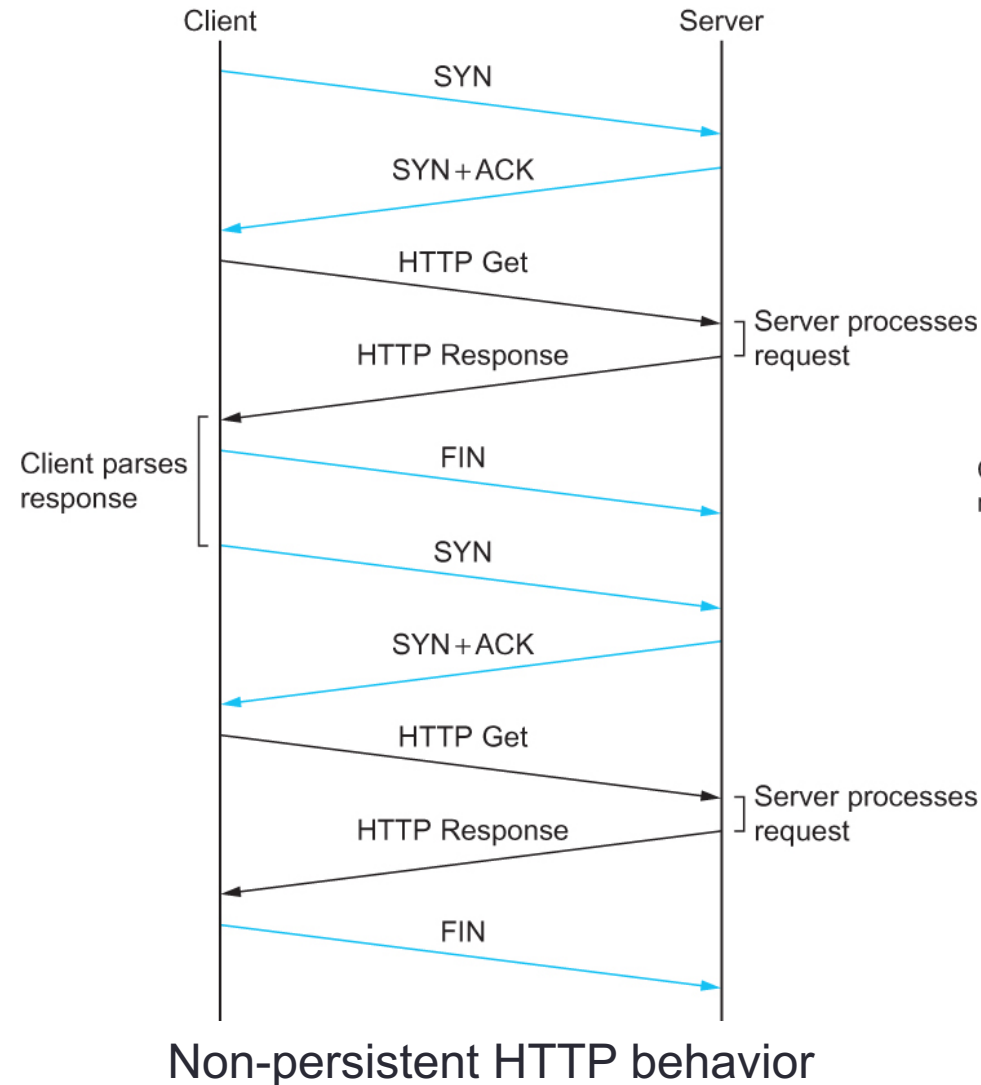


# Persistent HTTP

- **Non-persistent HTTP issues:**
  - requires at least 2 RTTs per object
  - OS must work and allocate host resources for each TCP connection
- **Persistent HTTP**
  - server leaves connection open after sending response
  - subsequent HTTP messages between same client/server are sent over connection



# Non-persistent HTTP vs. persistent HTTP



# Persistent HTTP

- Persistent without pipelining:
  - client issues new request only when previous response has been received
  - one RTT for each referenced object
- Persistent with pipelining:
  - default in HTTP/1.1
  - client sends requests as soon as it encounters a referenced object
  - as little as one RTT for all the referenced objects

# HTTP message: request

- Two types of HTTP messages: **request**, **response**
- HTTP request message:
  - ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

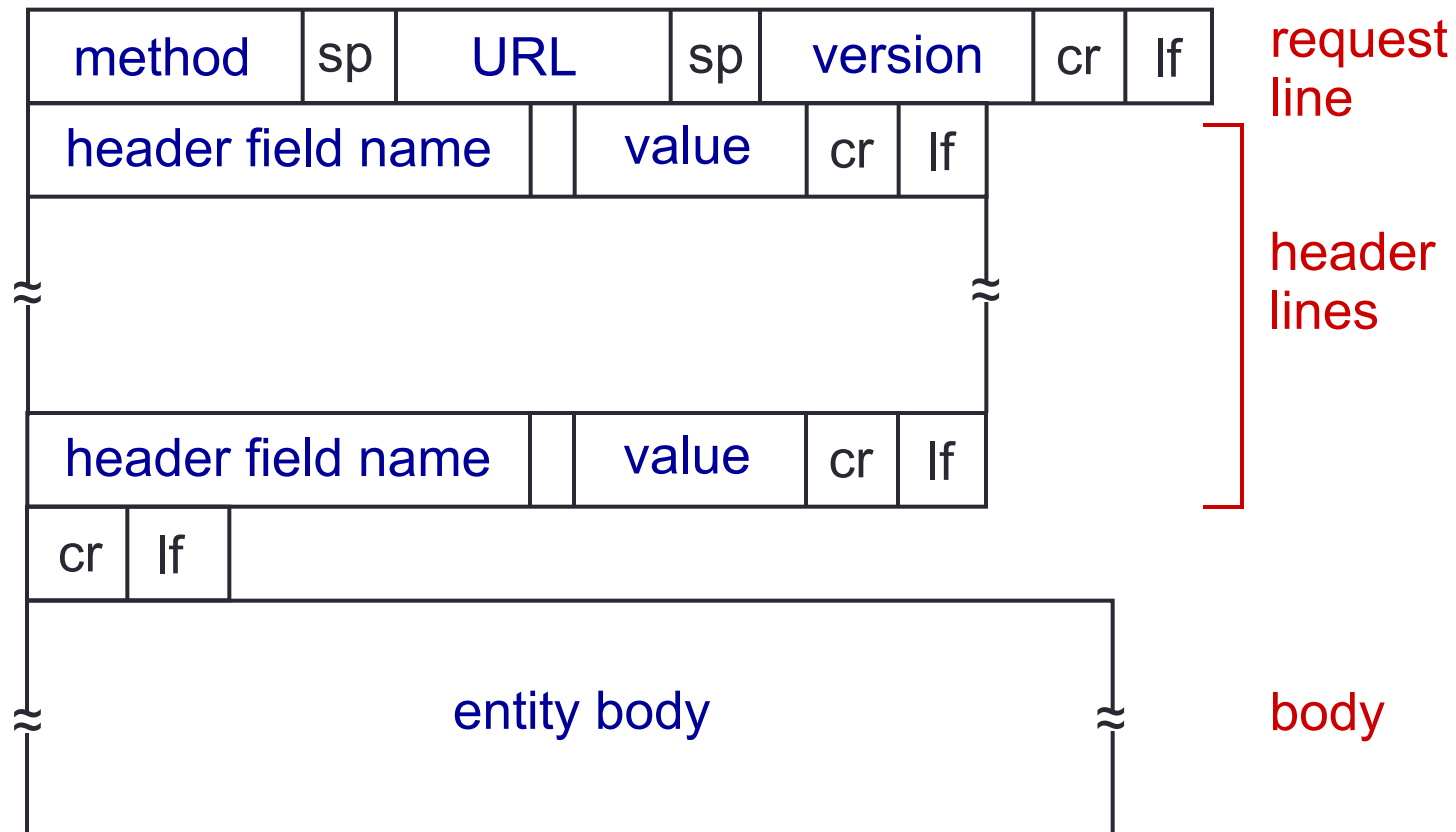
header  
lines

carriage return,  
line feed at start  
of line indicates  
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character  
line-feed character

# HTTP request: general format



sp: space

cr: carriage return (\r)

lf: line feed (\n)

# Uploading form input

- **POST** method:
  - Web page often includes form input
  - Input is uploaded to server in entity body
- **URL** method:
  - Uses GET method
  - Input is uploaded in URL field of request line:

**`www.somesite.com/animalsearch?monkeys&banana`**

# Method types

- HTTP/1.0
  - GET
  - POST
  - HEAD
    - asks server to leave requested object out of response
- HTTP/1.1
  - GET, POST, HEAD
  - PUT
    - uploads file in entity body to path specified in URL field
  - DELETE
    - deletes file specified in the URL field

# HTTP message: response

status line

(protocol

status code

status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

# HTTP response status codes

Status codes appear in the 1<sup>st</sup> line in server to client response message.

A few sample codes:

- **200 OK**
  - request succeeded, requested object later in this message
- **301 Moved Permanently**
  - requested object moved, new location specified later in this message (Location:)
- **400 Bad Request**
  - request message not understood by server
- **404 Not Found**
  - requested document not found on this server
- **505 HTTP Version Not Supported**



# Trying out HTTP

## 1. Telnet to a Web server:

```
telnet cs.binghamton.edu 80
```

Opens TCP connection to port 80 (default HTTP server port) at cs.binghamton.edu. Anything typed is sent to port 80 at cs.binghamton.edu

## 2. Type in an HTTP GET request:

```
GET /~yaoliu/courses/cs557/test.html HTTP/1.1  
Host: cs.binghamton.edu
```

By typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

## 3. Look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)

# Reading

- Textbooks on computer networks
  - Chapters on IP, TCP, and HTTP
- Socket tutorial by Paul Krzyzanowski@Rutgers
  - <http://www.cs.rutgers.edu/~pxk/352/notes/sockets/index.html>
- Beej's famous socket tutorial
  - <http://www.beej.us/guide/bgnet/html/multi/index.html>
- Socket programming documentation pages of your chosen programming language: Python (preferred), C/C++, Java.