


Physical and Logical Clocks

Yao Liu


Each machine has its own clock

- An event that occurred after another event may nevertheless be assigned an earlier time.



Perfect, fast, and slow clocks


- The relation between clock time and UTC when clocks tick at different rates.



Clock synchronization

- Physical clocks drift, therefore need for clock synchronization algorithms
 - Clock synchronization algorithms – Cristian's algorithm, NTP, Berkeley algorithm, etc.
- However, since we cannot perfectly synchronize clocks across computers, we cannot use physical time to order events

Clock synchronization using a time server



Cristian's algorithm

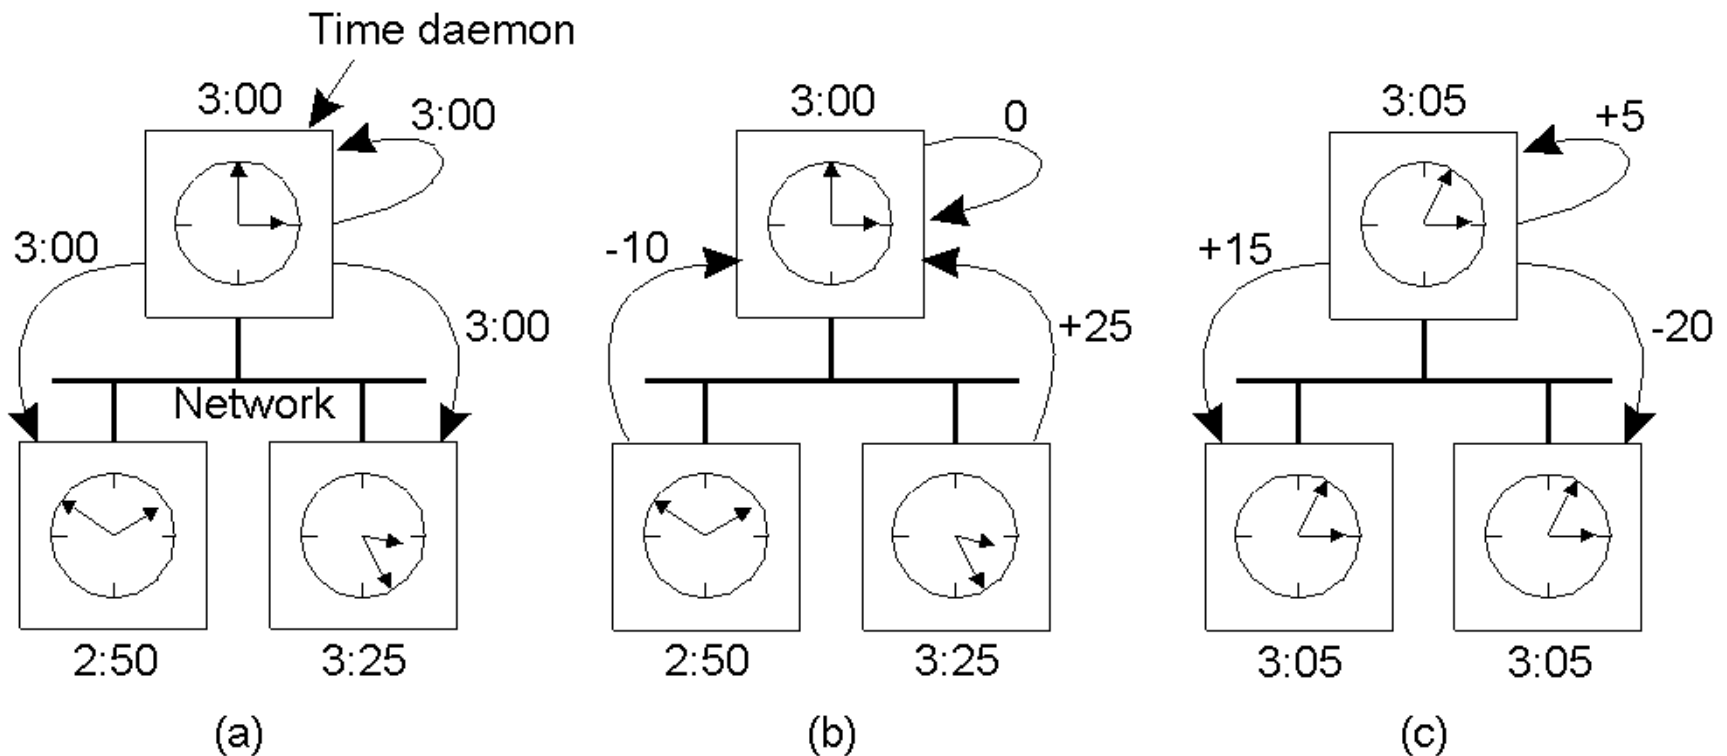
- A process p
 - requests the time in a message m_r
 - receives the time value t in a message m_t
- p should set its time to $t + T_{\text{round}}/2$

Cristian's algorithm


- If the value of the minimum transmission time min is known or can be conservatively estimated
 - Earliest time at which S could have placed its time in m_t was min after p dispatched m_r
 - Latest point at which it could do so was min before m_t arrived at p
 - Time by S's clock when message arrives at p is in range $[t + min, t + T_{round} - min]$
 - Accuracy $\pm(T_{round}/2 - min)$
- Further accuracy can be gained by making multiple requests to S and using the response with the shortest T_{round} .

The Berkeley algorithm

- a) The time daemon asks all the other machines for their clock values
- b) The machines answer
- c) The time daemon tells everyone how to adjust their clock



Network Time Protocol



Note: Arrows denote synchronization control, numbers denote strata.


Real synchronization is not perfect

- Clocks never exactly synchronized
- Often inadequate for distributed systems
 - might need to find out the order of any arbitrary pair of events
 - might need millionth-of-a-second precision

Logical time & clocks

- Lamport proposed using logical clocks based upon the “happened before” relation
 - If two events occur at the same process, then they occurred in the order observed
 - Whenever a message is sent between processes, the event of sending occurred before the event of receiving
 - X happened before Y denoted by $X \rightarrow Y$

Events occurring at three processes



Concurrency

- \rightarrow is only a partial-order
- Some events are unrelated
- We say e is concurrent with e' ($e \parallel e'$) if neither $e \rightarrow e'$ nor $e' \rightarrow e$


Lamport logical clocks

- Lamport clock orders events consistent with logical “happens before” ordering
- If $e \rightarrow e'$, then $L(e) < L(e')$

Lamport's algorithm

- Each process has its own logical clock
- Three rules:
 - C_p is incremented before each event at process p
 - When process p sends a message it piggybacks on it the value C_p
 - On receiving a message (m, t) , a process q computes $C_q = \max(C_q, t)$ and then applies the first rule before timestamping the receive event

Lamport timestamps for the events



Totally ordered logical clocks

- Lamport logical clocks only impose partial ordering
- For total order, use (T_a, P_a) where P_a is process id
- $(T_a, P_a) < (T_b, P_b)$ if and only if either $T_a < T_b$ or $(T_a = T_b \text{ and } P_a < P_b)$
- This ordering has no physical significance, but it is sometime useful, e.g., to break a tie between two processes trying to enter a critical section


Vector timestamps

- Shortcoming of Lamport's clocks:
 - if $L(e) < L(f)$, we **cannot** conclude that $e \rightarrow f$
- Vector timestamps
 - A process keeps a vector of integer clocks, one for each process
 - Like Lamport timestamps, processes piggyback vector timestamps on messages they send each other

Assigning vector timestamps

- Initially, $V_i[j] = 0$, for $i, j = 1, 2, \dots, N$.
- Just before process i timestamps an event, it sets
 - $V_i[i] = V_i[i] + 1$
- Process i includes the value $t = V_i$ in every message it sends.
- On receiving a message at process i :
 - $V_i[i] = V_i[i] + 1$
 - $V_i[j] = \max(V_i[j], t[j])$ for $j \neq i$

Vector timestamps for the events



Inferring happened-before relationship

- $V = W$ iff $V[j] = W[j]$ for $j = 1, 2, \dots, N$
- $V \leq W$ iff $V[j] \leq W[j]$ for $j = 1, 2, \dots, N$
- $V < W$ iff $V \leq W$ and $V \neq W$
 - $V \rightarrow W$, V happened before W
- If neither $V \leq W$ nor $W \leq V$
 - V and W are concurrent

Reading

- Sections 6.1 and 6.2 of TBook
- Sections 14.1-14.4 of CBook