


Snapshot Algorithm


Yao Liu

Detecting global properties


a. Garbage collection



b. Deadlock



c. Termination



Global snapshot

- Individual state of each **process** in the distributed system +
- Individual state of each **communication channel** in the distributed system
- Capture the instantaneous state of each process
- as well as the instantaneous state of each communication channel
 - i.e., messages in transit on the channels


If we have perfectly synchronized clocks

- Synchronize clocks of all processes
- Ask all processes to record their states at future time t
- However ...
 - There are no perfectly synchronized clocks
 - Messages in transit are not recorded
- In fact, synchronized is not required. Causality is enough.

Global states and consistent cuts

- How to capture a meaningful global state from local states recorded at different real times?
- Each process records events that correspond to internal actions, e.g., updating a variable, and the sending or receipt of a message
- A cut is a subset of the system's global history that is a union of prefixes of process histories
- A cut is **consistent** if for each event it contains, it also contains all events that happened before that event

Consistent and inconsistent cuts



Chandy and Lamport's snapshot algorithm

- Goal:
 - record a set of process and channel states for a set of processes
 - such that even if the combination of recorded states may never have occurred at the same time,
 - the recorded state is consistent
- State recorded locally at processes

Assumptions

- Neither channels nor processes fail; communication is reliable
- Channels are unidirectional and provide FIFO message delivery
- The graph of processes and channels is strongly connected (there is a path between any two processes)
- Any process may initiate a global snapshot at any time
- Processes may continue with their execution and send and receive normal messages while the snapshot takes place

Chandy and Lamport's snapshot algorithm

- For each process, P_i :
 - Incoming channels from which P_i receives messages
 - Outgoing channels to which P_i sends messages
- Each process must record both its local process state as well as states of all its incoming channels
- Marker messages
 - A special message used in the algorithm

Initiate the snapshot

- **P_i** initiates the snapshot by first recording its own local state
- for $j=1$ to N except i
 - **P_i sends** out a Marker message on outgoing channel **C_{ij}**
- **Starts recording** the incoming messages on each of the incoming channels at **P_i** : **C_{ji}** (for $j=1$ to N except i)

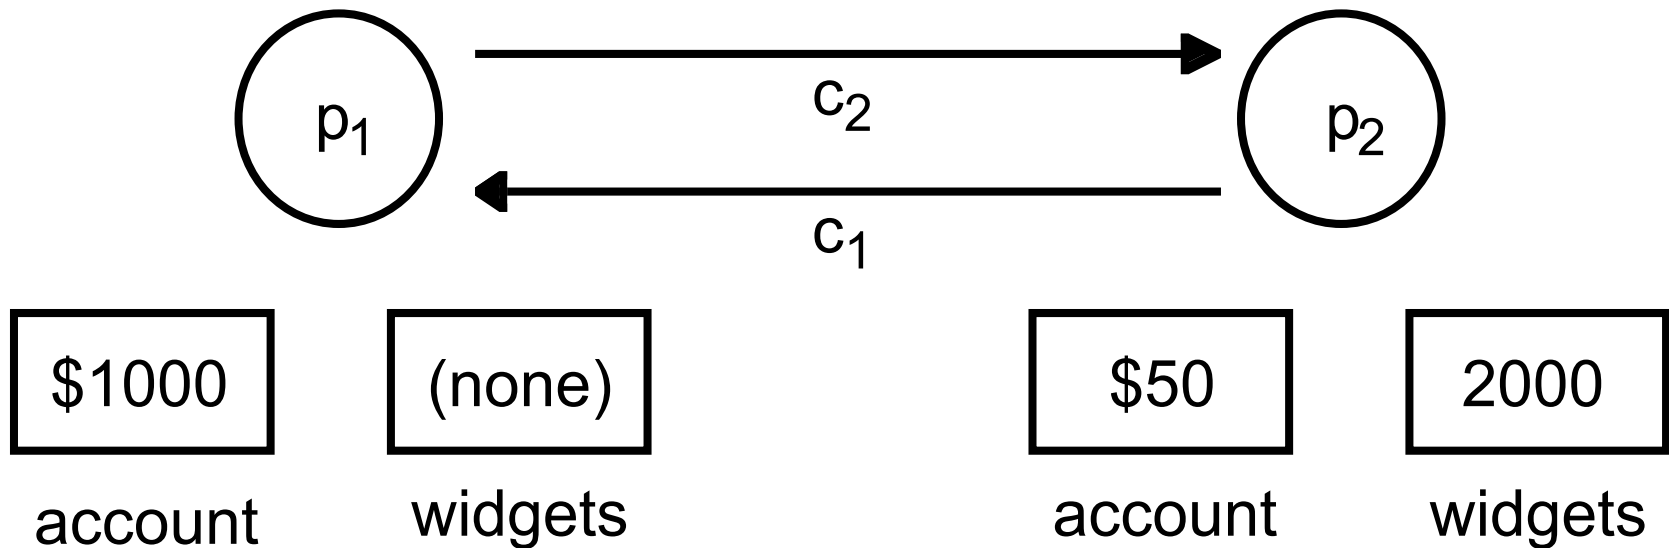
Marker receiving rule


Whenever a process P_i receives a Marker message on an incoming channel C_{ji} ($P_j \rightarrow P_i$)


- **if** (this is the first Marker P_i is seeing)
 - P_i records its own state first
 - Marks the state of channel C_{ji} as “empty”
 - for $j=1$ to N except i
 - P_i sends out a Marker message on outgoing channel C_{ij}
 - Starts recording the incoming messages on each of the incoming channels at P_i : C_{ji} (for $j=1$ to N except i)
- **else // already seen a Marker message**
 - Mark the state of channel C_{ji} as all the messages that have arrived on it since recording was turned on for C_{ji}
 - Stop recording on channel C_{ji}


Two processes and their initial states


- Assume that p_2 has already received an order for five widgets, which it will shortly dispatch to p_1

















P1: $(\$1000, 0)$

P2: $(\$50, 1995)$

C12: empty

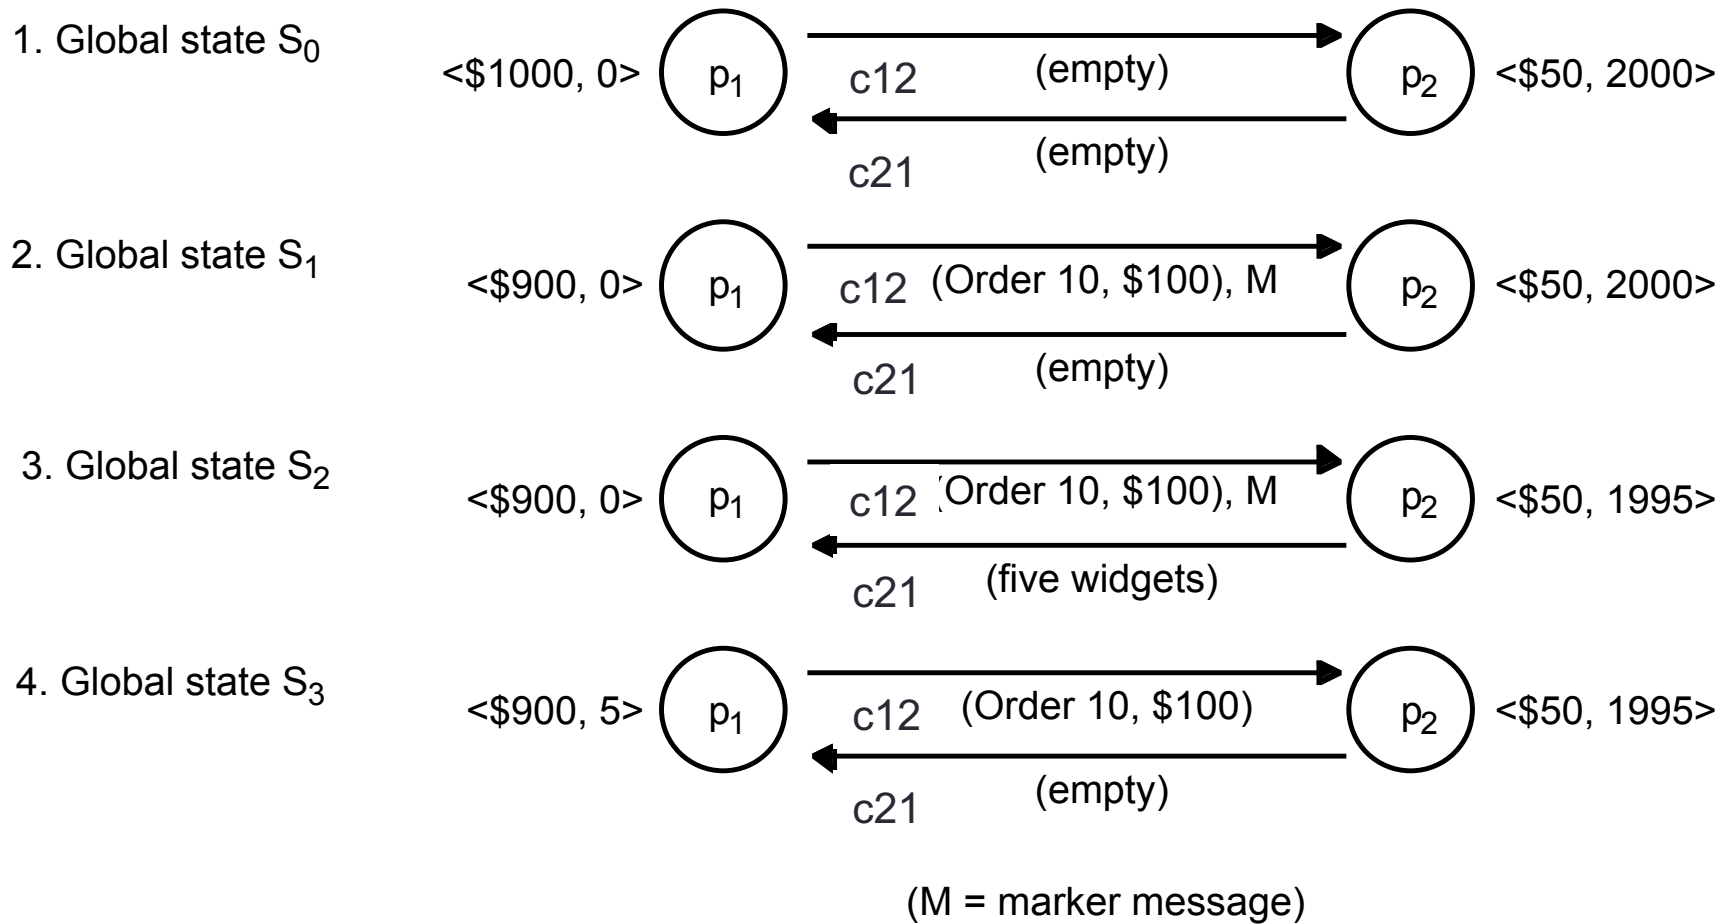
C21: (five widgets)



Marker Message



Regular Message



Final recorded state: P1: (\$1000,0), P2: (\$50, 1995), C12: empty, C21: (five widgets)

Reading

- Section 14.5 of CBook
- Paper on myCourses