

Key-value/NoSQL Store

Yao Liu

The Key-value abstraction

- (Business) Key \rightarrow Value
- (twitter.com) tweet id \rightarrow information about tweet
- (amazon.com) item number \rightarrow information about it
- (kayak.com) Flight number \rightarrow information about flight, e.g., availability
- (yourbank.com) Account number \rightarrow information about it

The Key-value abstraction

- It's a dictionary datastructure.
 - Insert, lookup, and delete by key
 - e.g., hash table, binary tree
- But distributed.
- Sound familiar? Remember Distributed Hash tables (DHT) in P2P systems?
- It's not surprising that key-value stores reuse many techniques from DHTs.

Database?

- Relational Database Management Systems (RDBMSs) have been around for ages
- MySQL is the most popular among them
- Data stored in tables
- Schema-based, i.e., structured tables
- Each row (data item) in a table has a primary key that is unique within that table
- Queried using SQL (Structured Query Language)
- Supports joins

Today's workloads

- Data: Large and unstructured
- Lots of random reads and writes
- Sometimes write-heavy
- Foreign keys rarely needed
- Joins infrequent

Needs of Today's Workloads

- Speed
- Avoid single point of failure
- Low TCO (total cost of operation)
- Fewer system administrators
- Incremental scalability
- Scale out, not scale up

Scale out, not Scale up

- Scale up = grow your cluster capacity by replacing with more powerful machines
 - Traditional approach
 - Not cost-effective, as you're buying above the sweet spot on the price curve
 - And you need to replace machines often
- Scale out = incrementally grow your cluster capacity by adding more COTS machines (Components Off the Shelf)
 - Cheaper
 - Over a long duration, phase in a few newer (faster) machines as you phase out a few older machines
 - Used by most companies who run datacenters and clouds today


Key-value/NoSQL data model

- NoSQL = “Not Only SQL”
- Necessary API operations:
 - `get(key)`
 - `put(key, value)`

CAP Theorem

- Proposed by Eric Brewer (Berkeley)
- Subsequently proved by Gilbert and Lynch (NUS and MIT)
- In a distributed system you can satisfy at most 2 out of the 3 guarantees:
 1. **Consistency**: all nodes see same data at any time, or reads return latest written value by any client
 2. **Availability**: the system allows operations all the time, and operations return quickly
 3. **Partition-tolerance**: the system continues to work in spite of network partitions

CAP Theorem



Why is Availability Important?

- Availability = Reads/writes complete reliably and quickly.
- Measurements have shown that a 500 ms increase in latency for operations at Amazon.com or at Google.com can cause a 20% drop in revenue.
- User cognitive drift: if more than a second elapses between clicking and material appearing, the user's mind is already somewhere else

Why is Consistency Important?


- Consistency = all nodes see same data at any time, or reads return latest written value by any client.
- When you access your bank or investment account via multiple clients (laptop, workstation, phone, tablet), you want the updates done from one client to be visible to other clients.
- When thousands of customers are looking to book a flight, all updates from any client (e.g., book a flight) should be accessible by other clients.

Why is Partition-Tolerance Important?

- Partitions can happen across datacenters when the Internet gets disconnected
 - Internet router outages
 - Under-sea cables cut
 - DNS not working
- Partitions can also occur within a datacenter, e.g., a rack switch outage
- Still desire system to continue functioning normally under this scenario

CAP Theorem Fallout

- Since partition-tolerance is essential in today's cloud computing systems,
- CAP theorem implies that a system has to choose between **consistency** and **availability**



Dynamo: Amazon's Highly Available Key-value Store

Amazon's Dynamo

- Distributed key-value storage
 - Only accessible with the primary key
 - `put(key, value)` & `get(key)`
- Used for many Amazon services
 - Shopping cart, best seller lists, customer preferences, etc.
- Inspired many NoSQL implementations
 - Cassandra
 - Riak
 - Project Voldemort

Design consideration

- Sacrifice strong consistency for availability
- “Always writeable” data store where no updates are rejected due to failures or concurrent writes.
- Conflict resolution is executed during ***read*** instead of ***write***, i.e., “always writeable”.

Design consideration (cont'd)

- Incremental scalability
 - be able to scale out one storage host at a time, with minimal impact on both operators of the system and the system itself.
- Symmetry
 - every node in Dynamo should have the same set of responsibilities as its peers.
- Decentralization
 - in the past, centralized control has resulted in outages and the goal is to avoid it as much as possible.
- Heterogeneity
 - this is essential in adding new nodes with higher capacity without having to upgrade all hosts at once.

System architecture

- Partitioning
- High Availability for writes
- Handling temporary failures
- Recovering from permanent failures
- Membership and failure detection

Partition algorithm

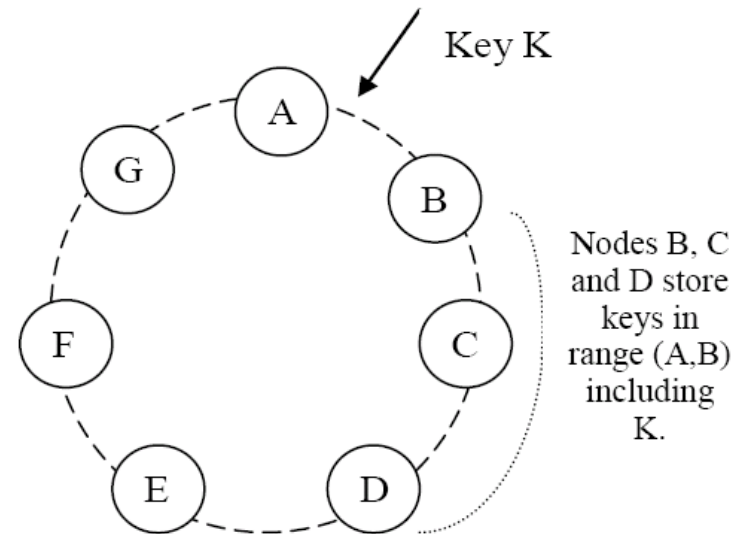
- Consistent hashing
 - the output range of a hash function is treated as a fixed circular space or “ring”.
 - With original consistent hashing, e.g., used by Chord DHT, load may become uneven
- “Virtual Nodes” for better load balancing
 - Each node can be responsible for more than one virtual nodes.

Partition algorithm (cont'd)

- Advantages of **using virtual nodes**:
 - If a node becomes unavailable, the load handled by this node is **evenly dispersed** across the remaining available nodes.
 - When a node becomes available again, or a new node is added to the system, the newly available node accepts **a roughly equivalent amount** of load from each of the other available nodes.
 - The number of virtual nodes that a node is responsible can be decided based on its **capacity**, accounting for **heterogeneity** in the physical infrastructure.

Replication

- Each data item is replicated at N hosts; configurable.
- The first is stored regularly with consistent hashing
- $N-1$ replicas are stored in the $N-1$ successor nodes
- “*preference list*”: The list of nodes that is responsible for storing a particular key.



Quorum

- Parameters:
 - N replicas
 - R readers
 - W writers
- Static quorum approach:
 - $R+W > N$
- Typical Dynamo configuration:
 - $N=3, R=2, W=2$

Sloppy quorum

- But we know that quorum-like approaches sacrifice availability for consistency!
- Dynamo uses a “sloppy” quorum.
 - all read and write operations are performed on the first N healthy nodes from the preference list, which may not always be the first N nodes encountered while walking the ring

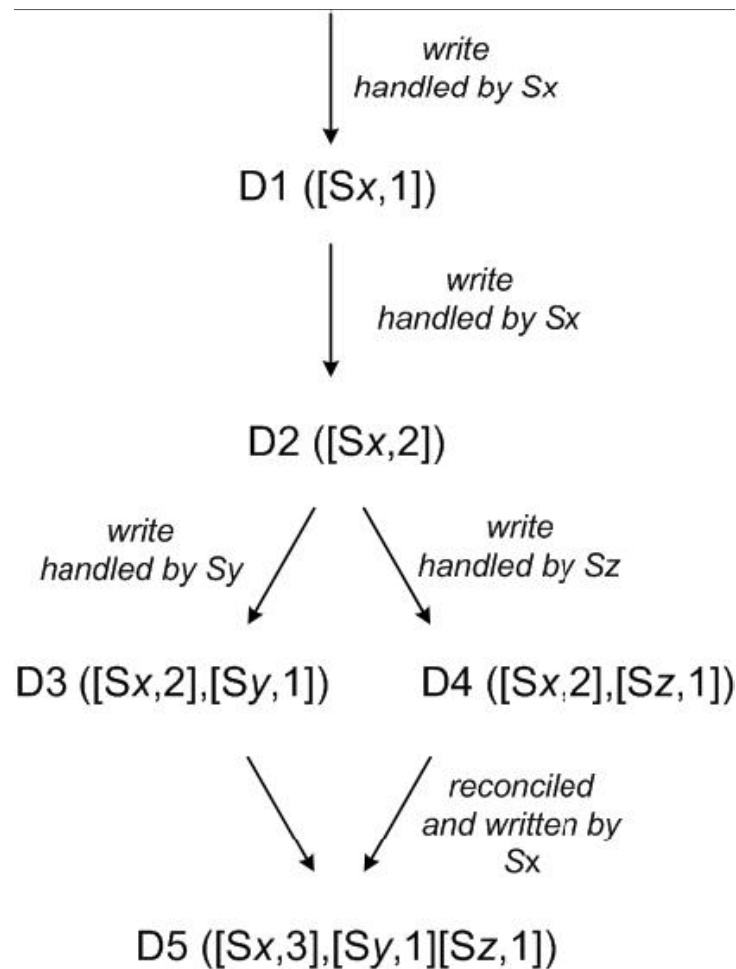
Data versioning

- Writes should succeed all the time
 - e.g., “Add to Cart”
- Need to reconcile inconsistent data due to network partitioning/failures
- Each object has a vector clock
 - e.g., D1 ([Sx, 1], [Sy, 1]): Object D1 has written once by server Sx and Sy, respectively.
 - Each node keeps all versions until the data becomes consistent
- If inconsistent, reconcile later.
 - e.g., deleted items might reappear in the shopping cart.

Vector clock

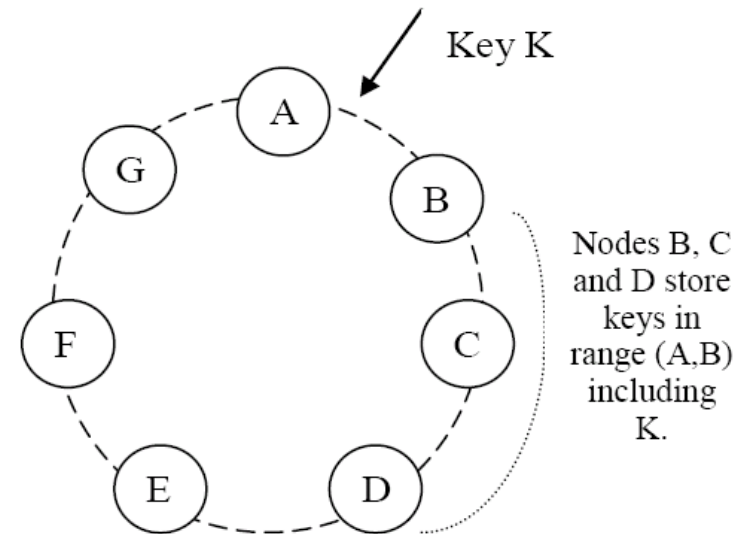
- A vector clock is a list of (node, counter) pairs.
- Every version of every object is associated with one vector clock.
- *If the counters on the first object's clock are less-than-or-equal to all of the nodes in the second clock, then the first is an ancestor of the second and can be forgotten.*

Vector clock example



Handling failures: hinted handoff

- Assume $N = 3$. When C is temporarily down or unreachable during a write, send replica to E.
- E is hinted that the replica belongs to C and it will deliver to C when C is recovered.
- Again: “always writeable”




Replica synchronization

- Key ranges are replicated.
- If a node fails and recovers, it needs to quickly determine whether it needs to resynchronize or not.
 - Transferring entire (key, value) pairs for comparison is not an option.
- Merkle tree:
 - a hash tree where leaves are hashes of the values of individual keys.
 - Parent nodes higher in the tree are hashes of their respective children.

Replica synchronization

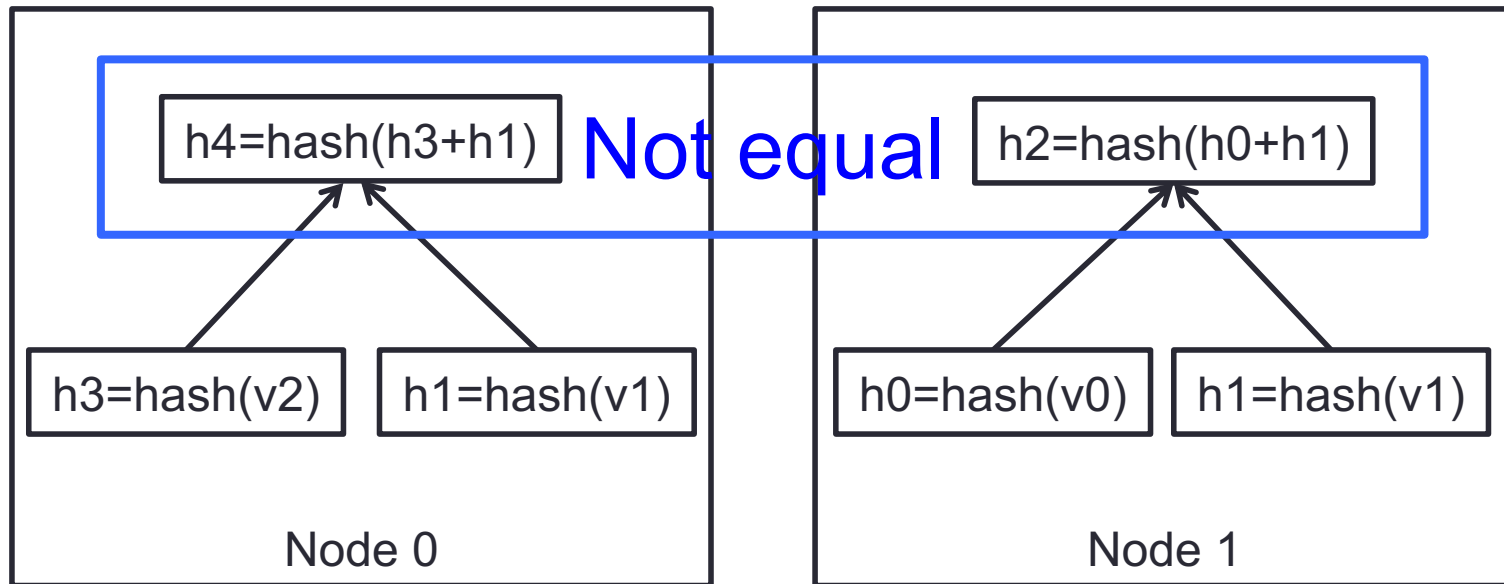
- Comparing two nodes that are **synchronized**
 - Two (key, value) pairs: (k_0, v_0) & (k_1, v_1)



+ indicates concatenate

Replica synchronization

- Comparing two nodes that are *not synchronized*
 - One: (k_0, v_2) & (k_1, v_1)
 - The other: (k_0, v_0) & (k_1, v_1)



+ indicates concatenate

Membership

- Nodes are organized as a ring just like Chord using consistent hashing
- But everyone knows everyone else.
- Node join/leave
 - Manually done
 - An operator uses a console to add/delete a node
 - Reason: it's a well-maintained system; nodes come back pretty quickly and don't depart permanently most of the time
- Membership change propagation
 - Each node maintains its own view of the membership & the history of the membership changes
 - Gossip-based protocol for propagation (each node contacts a randomly selected node every second)
- Eventually-consistent membership protocol

Failure detection

- Avoid communicating with unreachable peers
 - during put() and get()
 - transferring partitions and hinted replicas
- Does not use a separate protocol; each request serves as a ping
 - Dynamo has enough requests at any moment anyway
- If a node doesn't respond to a request, it is considered to be failed.

Conclusion

- Amazon Dynamo is a **highly available** and **scalable** data store.
- Techniques:
 - Gossiping for propagation of membership changes
 - Consistent hashing for node and key distribution
 - Data versioning for eventually-consistent data objects
 - Sloppy quorum for partition/failure tolerance
 - Merkle tree for replica synchronization

Reading

- Dynamo: Amazon's Highly Available Key-value Store
- <http://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>

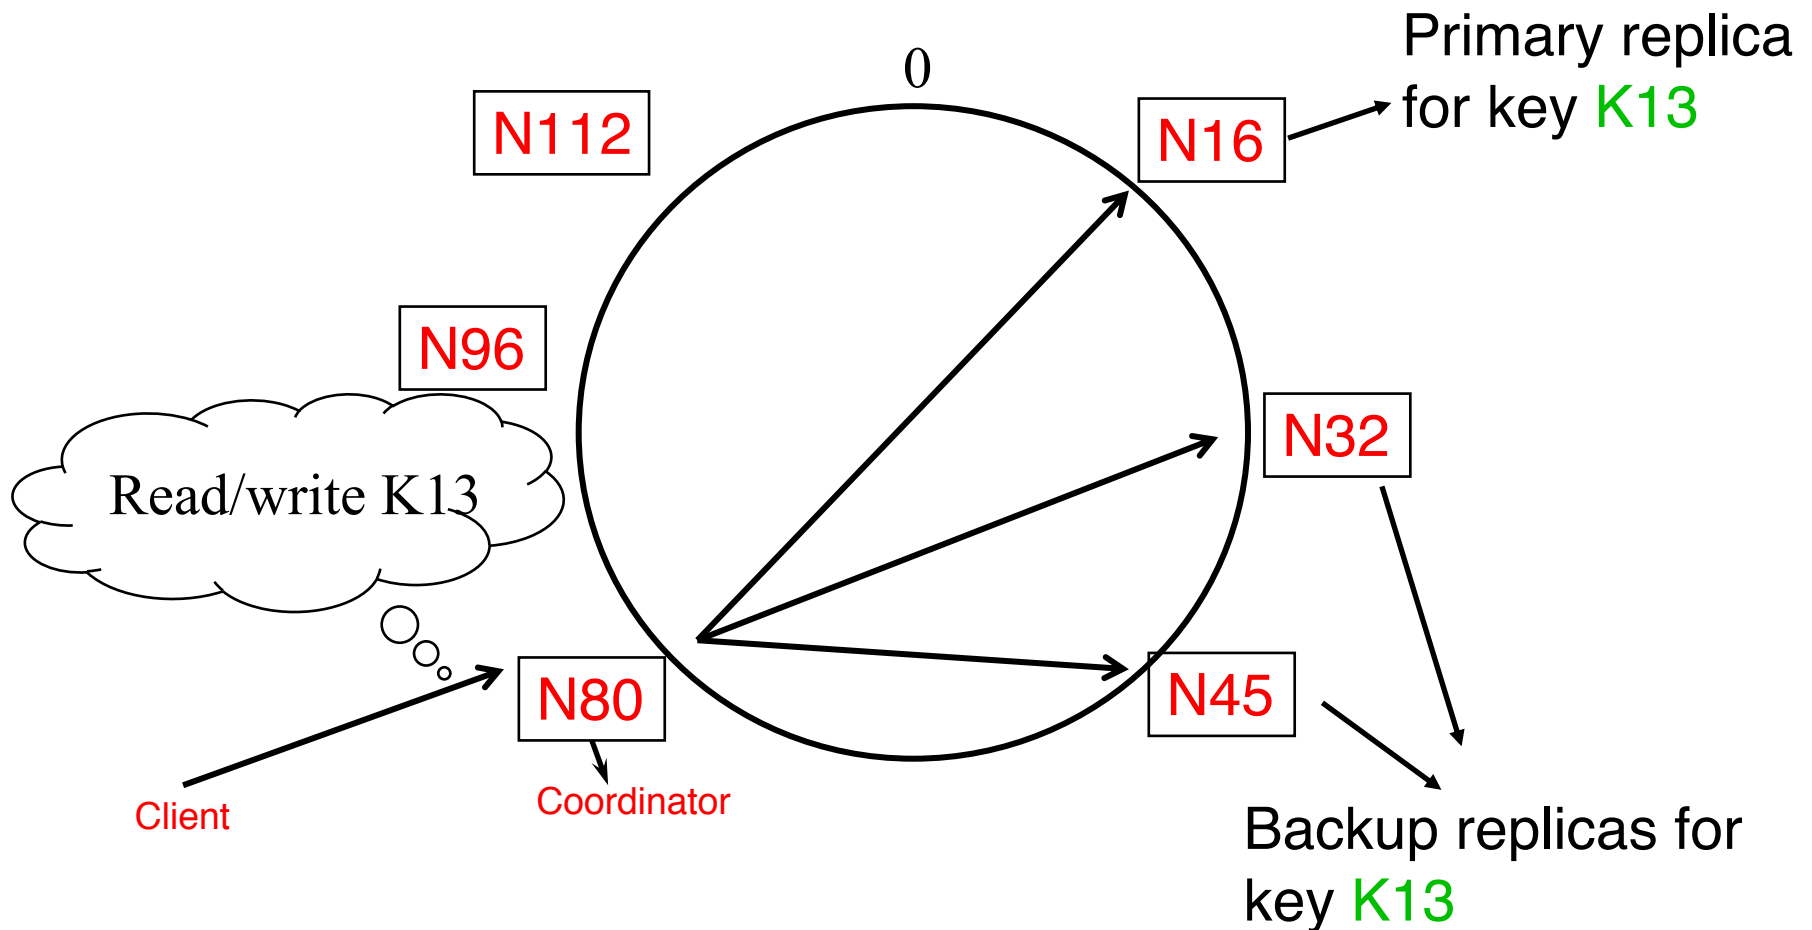
Cassandra

Cassandra

- A distributed key-value store
- Intended to run in a datacenter (and also across DCs)
- Originally designed at Facebook
- Open-sourced later, today an Apache project
- Some of the companies that use Cassandra in their production clusters
 - IBM, Adobe, HP, eBay, Ericsson, Symantec
 - Twitter, Spotify
 - Netflix: uses Cassandra to keep track of your current position in the video you're watching

Partitioning

- Which server(s) a key-value resides on?



Partitioner


- A function for computing the key \rightarrow server mapping

Data replication strategies

- Two replication strategies are available:
 - `SimpleStrategy`
 - `NetworkTopologyStrategy`
- **SimpleStrategy**: uses `Partitioner`
 - *RandomPartitioner*: Chord-like hash partitioning
 - *ByteOrderedPartitioner*: Assigns ranges of keys to servers.
 - Easier for range queries (e.g., return all twitter users starting with [a-b])
 - places the first replica on a node determined by the partitioner. Additional replicas are placed on the next nodes clockwise in the ring without considering topology (rack or datacenter location)
- **NetworkTopologyStrategy**: for multi-datacenter deployments
 - Two replicas per DC
 - Three replicas per DC
 - Within DC: first replica placed according to `Partitioner`, then go clockwise around ring until reaching the first node in another rack

SimpleStrategy

Replication factor: 3




NetworkTopologyStrategy

Replication factor: {DC1:2, DC2:2}

Node	DC	Rack
N0	DC1	Rack2
N42	DC1	Rack1
N108	DC1	Rack1
N20	DC2	Rack1
N64	DC2	Rack1
N85	DC2	Rack2

Key10 replicas: DC1:{N42,N0}, DC2:{N20,N85}



Snitches

- Maps: IPs to racks and DCs. Configured in `cassandra.yaml` config file
- Some options:
 - `SimpleSnitch`: Unaware of topology (rack-unaware)
 - `RackInferring`: Assumes topology of network by octet of server's IP address
 - `101.201.202.203 = x.<DC octet>.<rack octet>.<node octet>`
 - `PropertyFileSnitch`: uses a config file
 - `EC2Snitch`: uses Amazon EC2.
 - EC2 Region = DC
 - Availability zone = rack
- Other snitch options available

Write

- Need to be lock-free and fast (no reads or disk seeks)
- Client sends write to one coordinator node in Cassandra cluster
- Coordinator uses Partitioner to send query to all replica nodes responsible for key
- When X replicas respond, coordinator returns an acknowledgement to the client
 - X ? We'll see later.

Write

- Always writable: Hinted Handoff mechanism
 - If any replica is down, the coordinator writes to all other replicas and keeps the write locally until down replica comes back up.
 - When all replicas are down, the Coordinator (front end) buffers writes (for up to a few hours).
- Multi-datacenter write
 - The coordinator forwards the write request to one replica in each of other datacenters, who will further forward the write to local replicas

Write on a replica node

On receiving a write

1. Log it in **disk commit log** (for failure recovery)
2. Make changes to appropriate memtables
 - **Memtable** = In-memory representation of multiple key-value pairs
 - *Typically append-only datastructure (fast)*
 - Cache that can be searched by key
 - Write-back cache as opposed to write-through

Later, when memtable is full or old, flush to disk

- Data File: An **SSTable** (Sorted String Table) – list of key-value pairs, sorted by key
- *SSTables are immutable (once created, they don't change)*
- Index file: An SSTable of (key, position in data SSTable) pairs
- And a Bloom filter (for efficient search)

Column-oriented storage

- RDBMSs store an entire row together (on disk or at a server)
- NoSQL systems typically store a column together (or a group of columns).
- Why useful?
 - Range searches within a column are fast since you don't need to fetch the entire database
 - e.g., get me all the `blog_ids` from the `blog` table that were updated within the past month
 - Search in the `last_updated` column, fetch corresponding `blog_id` column
 - Don't need to fetch the other columns

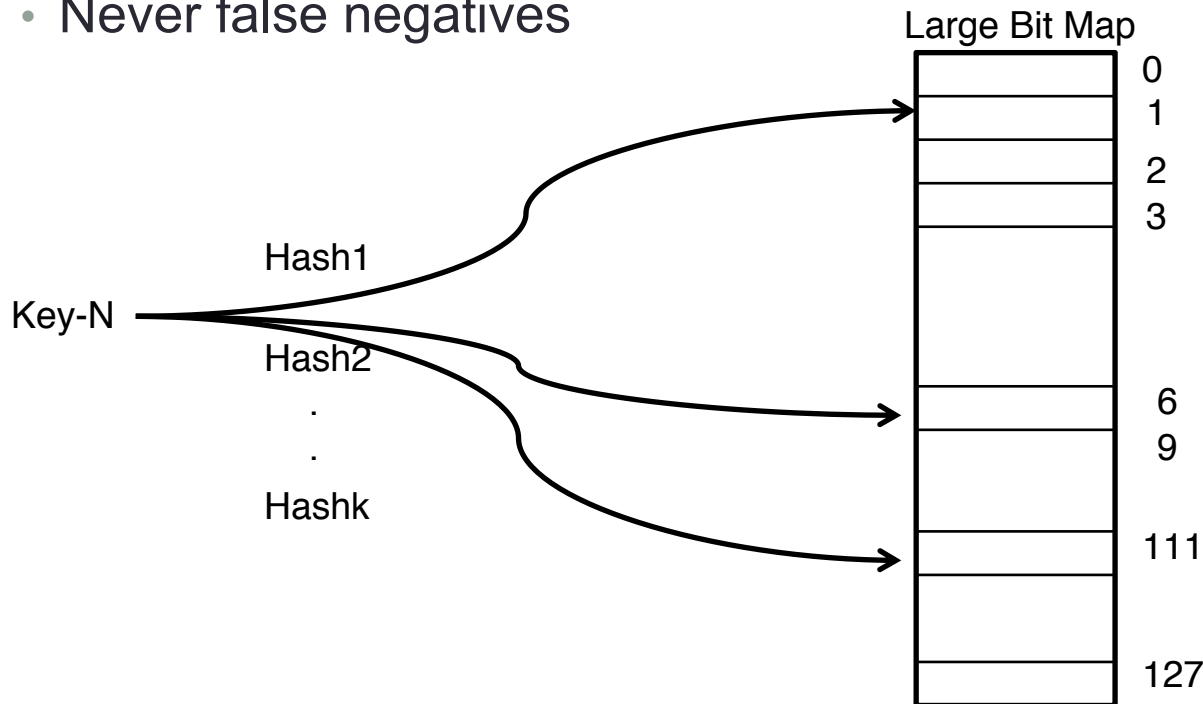
Compaction

Data updates accumulate over time and SSTables and logs need to be compacted

- The process of compaction merges SSTables, i.e., by merging updates for a key
- Run periodically and locally at each server

Bloom filter

- Compact way of representing a set of items
- Checking for existence in set is cheap
- Some probability of false positives
 - an item not in set may check true as being in set
- Never false negatives



On insert, set all hashed bits.

On check-if-present, return true if all hashed bits set.

- False positives

False positive rate low

- $k = 4$ hash functions
- $n = 100$ items
- $m = 3200$ bits
- FP rate = 0.0191%

Delete

Delete: don't delete item right away

- Add a **tombstone**
- Eventually, when compaction encounters tombstone it will delete item

Read

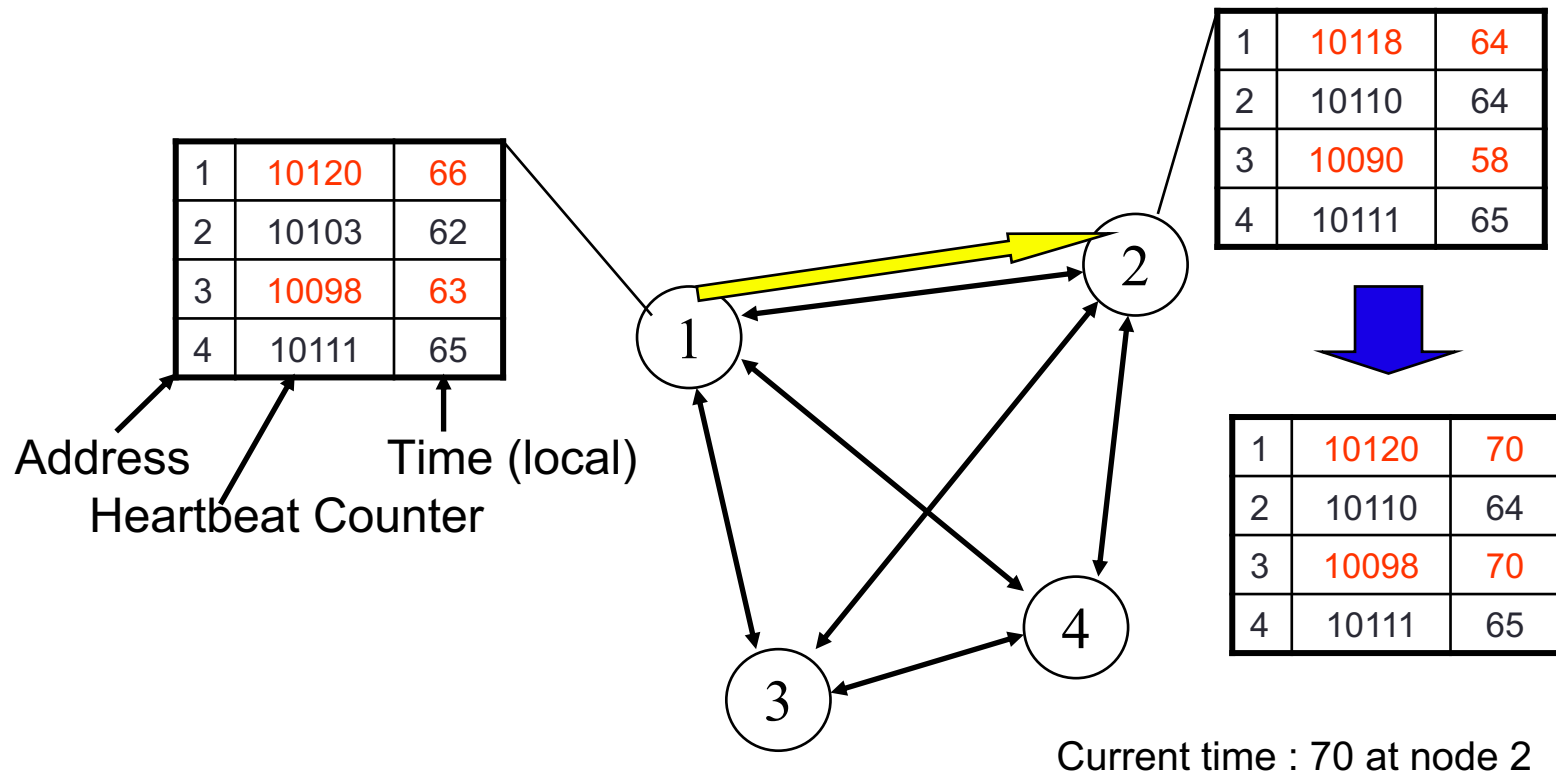
Read: Similar to writes, except

- Coordinator can contact X replicas
 - Coordinator sends read to replicas that have responded quickest in past
 - When X replicas respond, coordinator returns the **latest-timestamped value** from among those X
 - (X? We'll see later.)
- Coordinator also fetches value from other replicas
 - Checks consistency in the background, initiating a **read repair** if any two values are different
 - This mechanism seeks to eventually bring all replicas up to date
- At a replica
 - Read looks at **Memtables** first, and then **SSTables**
 - A row may be split across multiple SSTables → reads need to touch multiple SSTables → reads slower than writes (but still fast)

Membership

- Any server in cluster could be the coordinator
- So every server needs to maintain a list of all the other servers that are currently in the cluster
- List needs to be updated automatically as servers join, leave, and fail
- Gossip-based
 - Nodes periodically gossip their membership list
 - On receipt, the local membership list is updated
 - If any heartbeat older than T_{fail} , node is marked as failed

Gossip-based cluster membership



Failure detection

- Does not use a fixed threshold for marking failing nodes.
- Accrual failure detector
 - designed to adapt to changing network conditions
- The value output, PHI, represents a suspicion level.
- Applications set an appropriate threshold, trigger suspicions and perform appropriate actions.

Cassandra vs. RDBMS

- MySQL is one of the most popular RDBMS(and has been for a while)
- On > 50 GB data
- MySQL
 - Writes 300 ms avg
 - Reads 350 ms avg
- Cassandra
 - Writes 0.12 ms avg
 - Reads 15 ms avg
- Orders of magnitude faster

Consistency levels in Cassandra

- Client is allowed to choose a consistency level for each operation (read/write)
 - ANY: any server (may not be replica)
 - Fastest: coordinator caches write and replies quickly to client
 - ALL: all replicas
 - Ensures strong consistency, but slowest
 - ONE: at least one replica
 - Faster than ALL
 - QUORUM: quorum across all replicas in all datacenters (DCs)
 - Global consistency, but still fast
 - LOCAL_QUORUM: quorum in coordinator's DC
 - Faster: only waits for quorum in the first DC the client contacts
 - EACH_QUORUM: quorum in every DC
 - Lets each DC do its own quorum: supports hierarchical replies

Quorum in detail

- Reads
 - Client specifies value of R ($\leq N$ = total number of replicas of that key).
 - R = read consistency level.
 - Coordinator waits for R replicas to respond before sending result to client.
 - In background, coordinator checks for consistency of remaining $(N-R)$ replicas, and initiates read repair if needed.

Quorum in detail

- Writes come in two flavors
 - Client specifies W ($\leq N$)
 - W = write consistency level.
 - Client writes new value to W replicas and returns. Two flavors:
 - Coordinator blocks until quorum is reached.
 - Asynchronous: push the data to appropriate nodes but return immediately

Reading

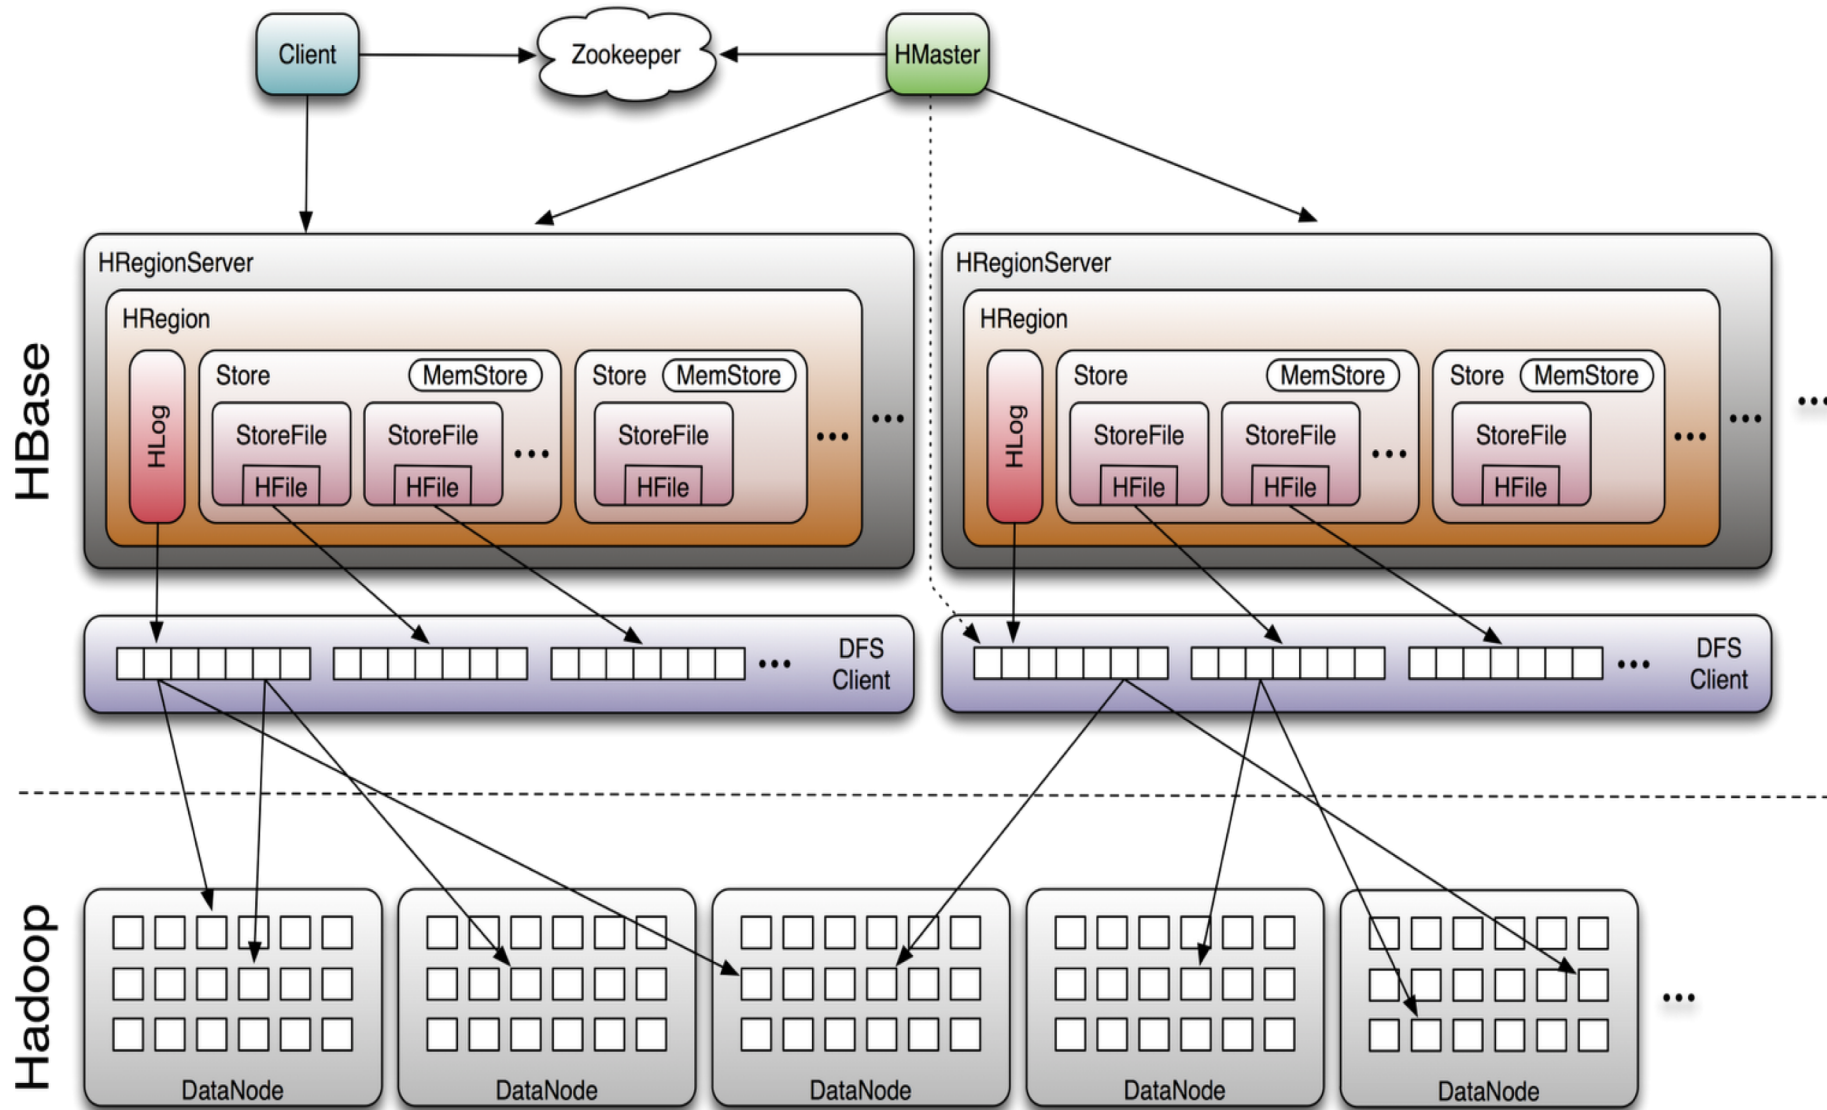
- Cassandra
- <https://www.cs.cornell.edu/projects/ladis2009/papers/lakshman-ladis2009.pdf>
- <http://www.slideshare.net/Eweaver/cassandra-presentation-at-nosql>
- <http://cassandra.apache.org/doc/latest/architecture/index.html>

HBase

HBase

- BigTable is a distributed system for storing structured data at Google
- Yahoo! Open-sourced it → HBase
- Major Apache project today
- Facebook uses HBase internally
- API functions
 - Get/Put(row)
 - Scan(row range, filter) – range queries
 - MultiAction for batch processing
- Unlike Cassandra, HBase prefers consistency (over availability)

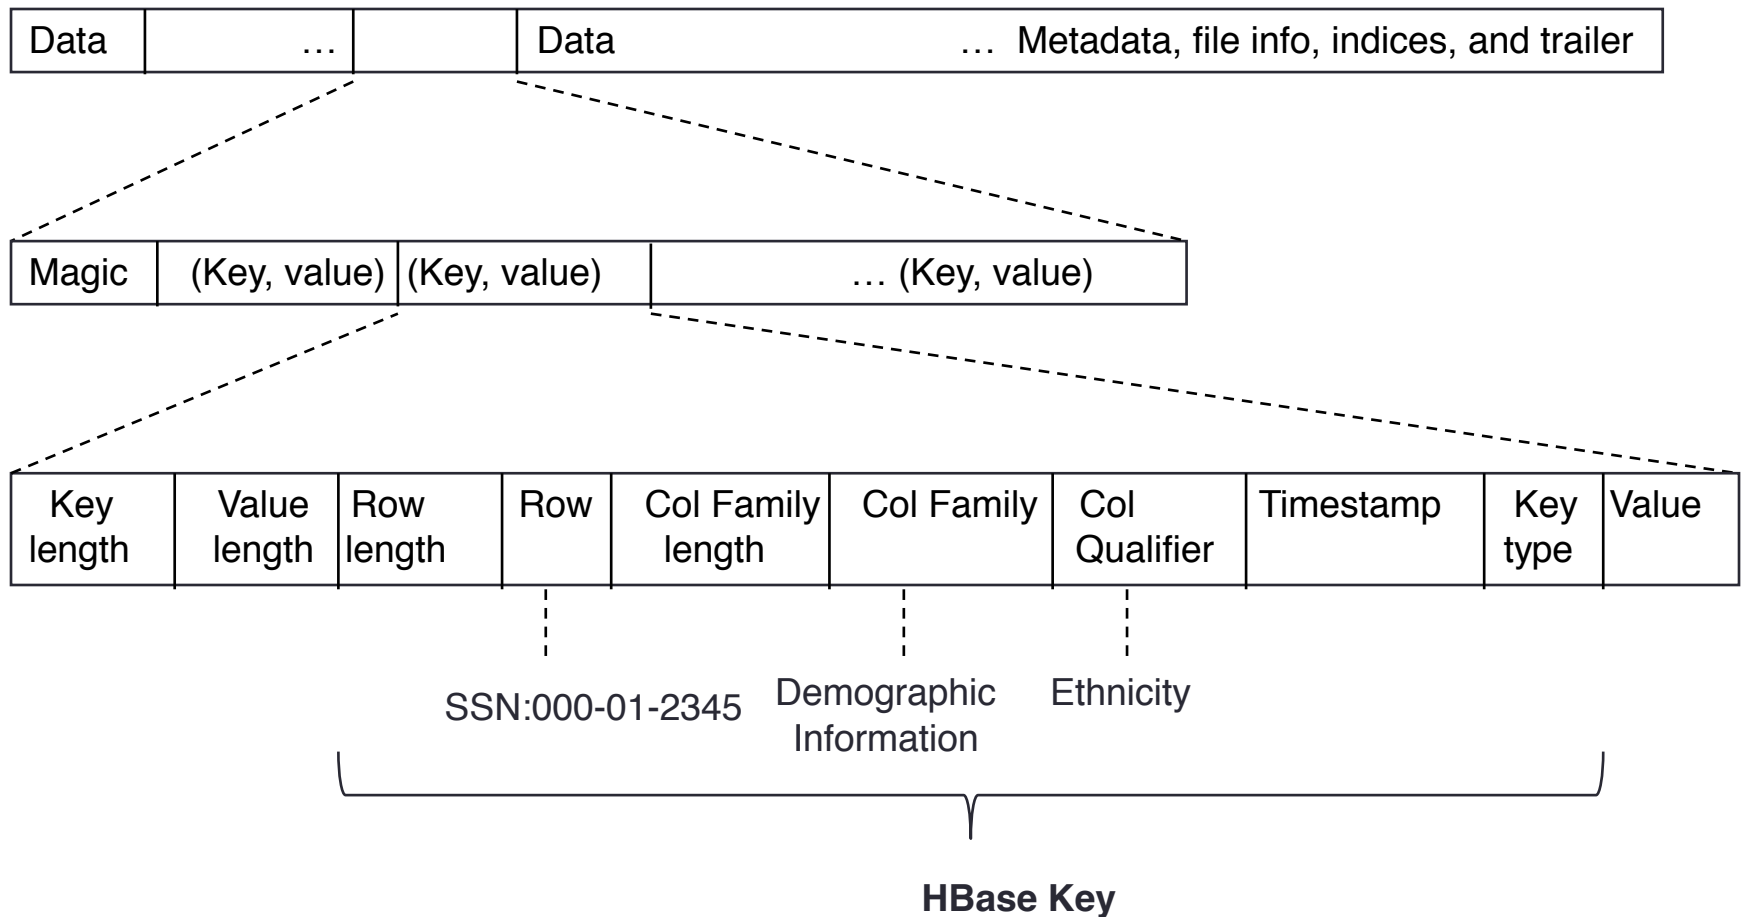
HBase architecture



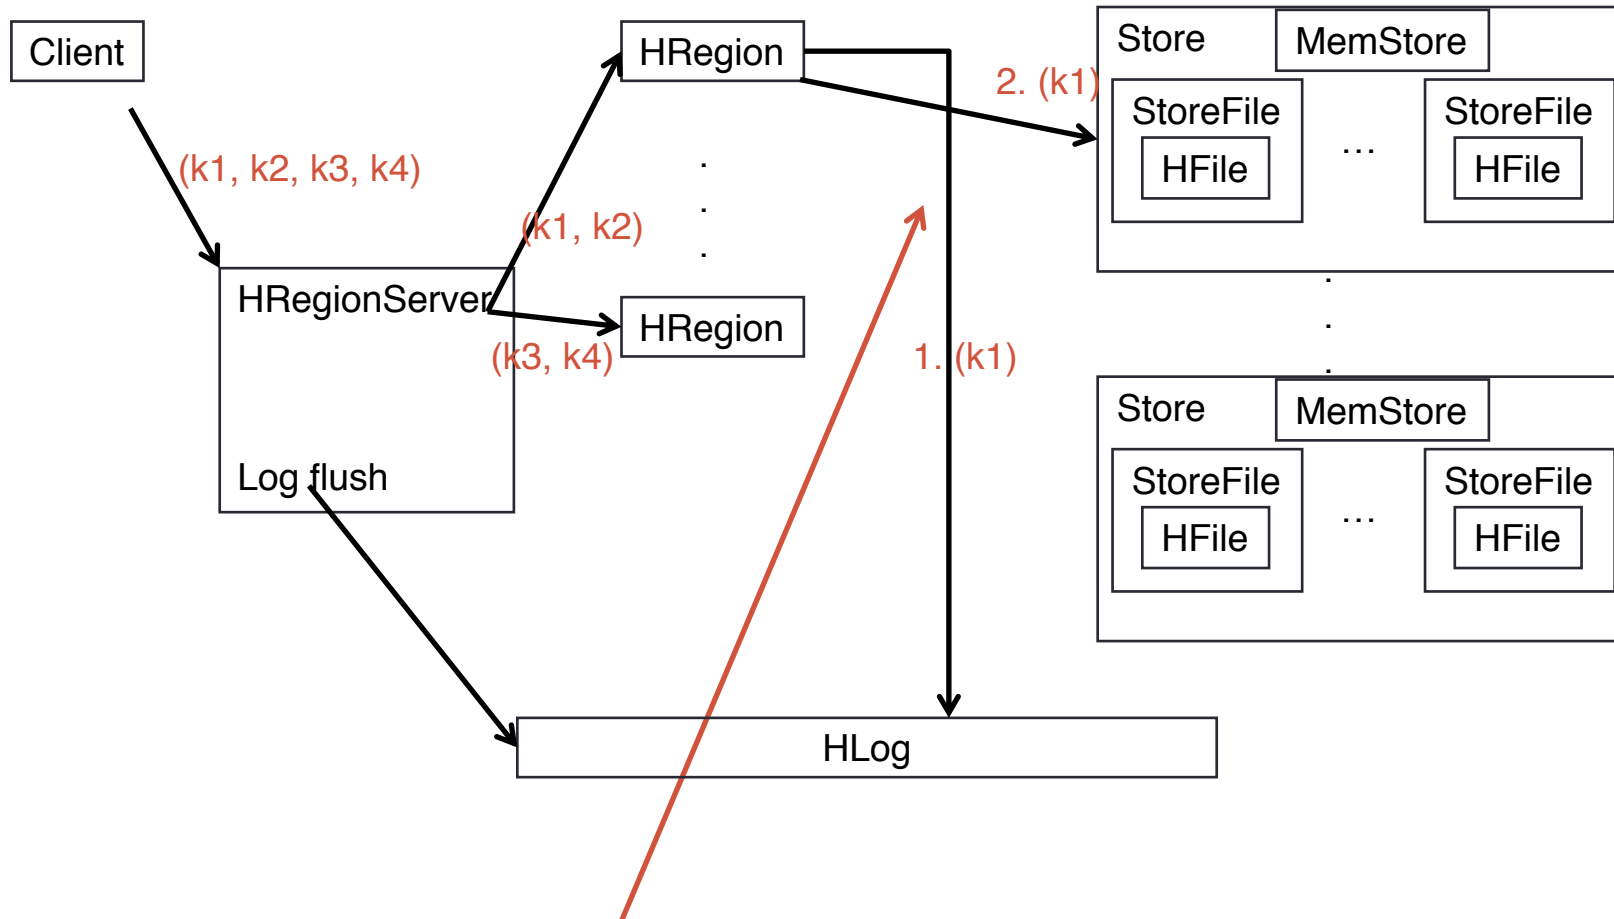
HBase storage hierarchy

- HBase Table
 - Split it into multiple regions:
 - a region is only served by one single regionserver
 - a row can only belong to one region
 - ColumnFamily = subset of columns with similar query patterns
 - One Store per combination of ColumnFamily + region
 - Memstore for each store: in-memory updates to Store; flushed to disk when full
 - StoreFiles for each store for each region: where the data lives
 - also called HFile
- HFile
 - SSTable from Google's BigTable

HFile

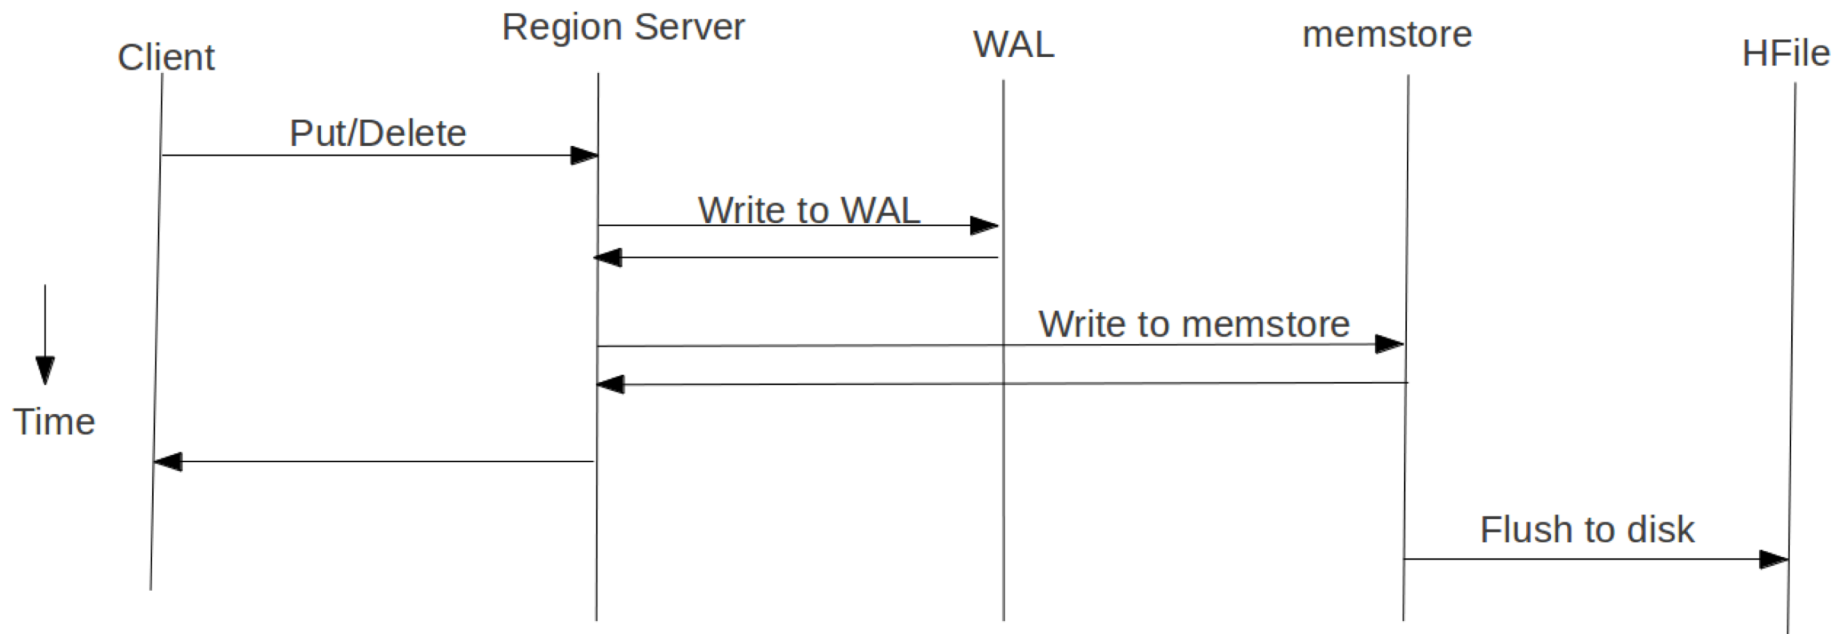


Strong consistency: HBase write-ahead log (WAL)



Write to HLog before writing to MemStore
Helps recover from failure by replaying HLog.

HBase write path



HBase Write Path

When a regionserver fails...

- The regions immediately become unavailable because the RegionServer is down.
- The Master will detect that the RegionServer has failed.
- The region assignments will be considered invalid and will be re-assigned just like boot up.

Log replay

- After recovery from failure, or upon boot up (HRegionServer/HMaster)
 - Replay any stale logs (use timestamps to find out where the database is w.r.t. the logs)
 - Replay: add edits to the MemStore

Reading

- BigTable:
- <http://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>
- HBase:
- <http://hbase.apache.org/book.html>

MongoDB

Data model

- Stores data in form of BSON (Binary JavaScript Object Notation) *documents*

```
{  
    name: "travis",  
    salary: 30000,  
    designation: "Computer Scientist",  
    teams: [ "front-end", "database" ]  
}
```


- Group of related *documents* with a shared common index is a *collection*

MongoDB: typical query

- Query all employee names with salary greater than 18000 sorted in ascending order

• `db.employee.find({salary:{$gt:18000}, {name:1}}).sort({salary:1})`

collection condition projection modifier



Insert

Insert a row entry for new employee Sally

```
db.employee.insert({  
    name: "sally",  
    salary: 15000,  
    designation: "MTS",  
    teams: [ "cluster-management" ]  
})
```

Update

All employees with salary greater than 18000 get a designation of Manager

<i>Update Criteria</i>	<code>{salary:{\$gt:18000}},</code>
<i>Update Action</i>	<code>{\$set: {designation: "Manager"}},</code>
<i>Update Option</i>	<code>{multi: true}</code>
	<code>)</code>

Multi-option allows multiple document update

Delete

Remove all employees who earn less than 10000

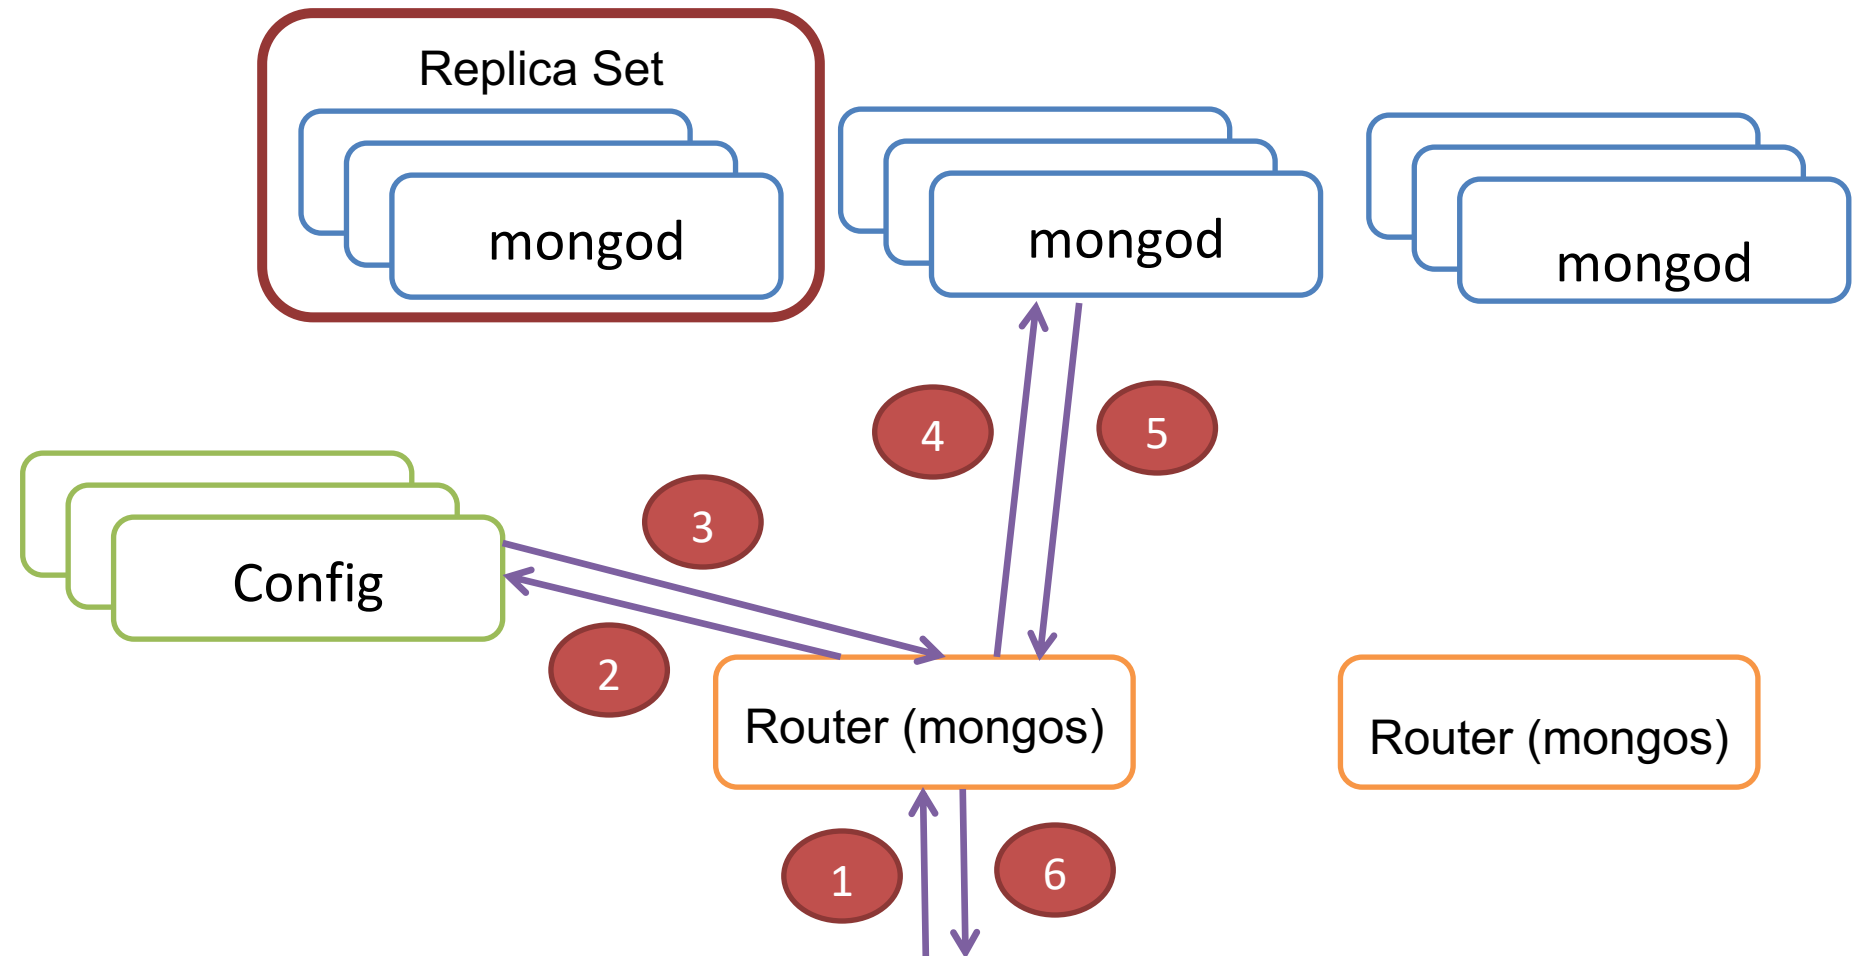
Remove Criteria `db.employee.remove(
 {salary:{<10000}},
)`

Can accept a flag to limit the number of documents removed

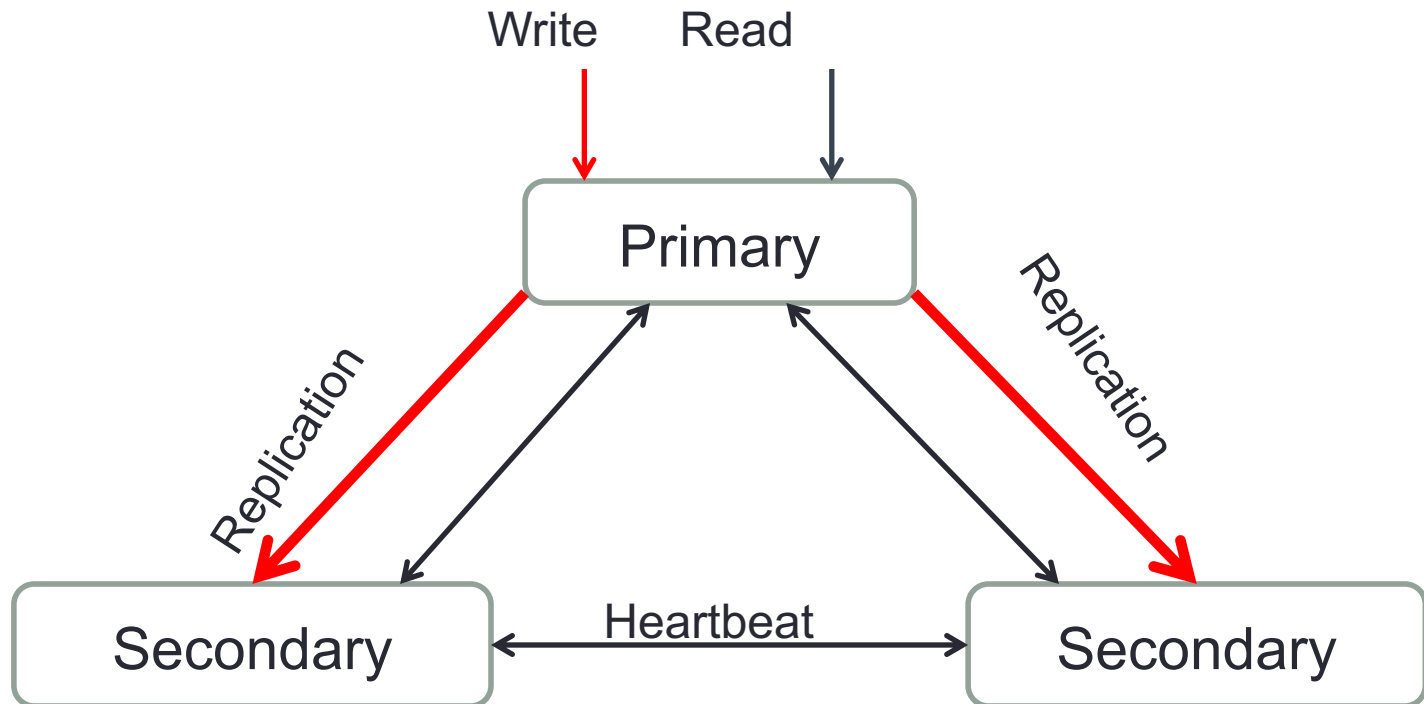
Typical MongoDB deployment

- Data split into **chunks**, based on shard key (~ primary key)
 - Either use hash or range-partitioning
- **Shard**: collection of chunks
- Shard assigned to a replica set
- **Replica set** consists of multiple mongod servers (typically 3 mongod's)
- Replica set members are mirrors of each other
 - One is primary
 - Others are secondaries
- **Routers**: **mongos** server receives client queries and routes them to right replica set
- **Config server**: stores collection level metadata

Typical MongoDB deployment



Replication



Replication

- Uses an oplog (operation log) for data sync
 - Oplog maintained at primary, delta transferred to secondary continuously/every once in a while
- When needed, leader election protocol elects a master
- Some mongod servers do not maintain data but can vote – called as Arbiters

Read preference

- Determine where to route read operation
- Default is primary. Some other options are
 - primary-preferred
 - secondary
 - nearest
- Helps reduce latency, improve throughput
- Reads from secondary may fetch stale data

Write concern

- Determines the guarantee that MongoDB provides on the success of a write operation
- Default is *acknowledged* (primary returns answer immediately).
 - Other options are
 - journaled (typically at primary)
 - replica-acknowledged (quorum with a value of W), etc.
- Weaker write concern implies faster write time

Write operation performance

- Journaling: Write-ahead logging to an on-disk journal for durability
- Indexing: Every write needs to update every index associated with the collection

Balancing

- Over time, some chunks may get larger than others
- Splitting: Upper bound on chunk size; when hit, chunk is split
- Balancing: Migrates chunks among shards if there is an uneven distribution

Consistency

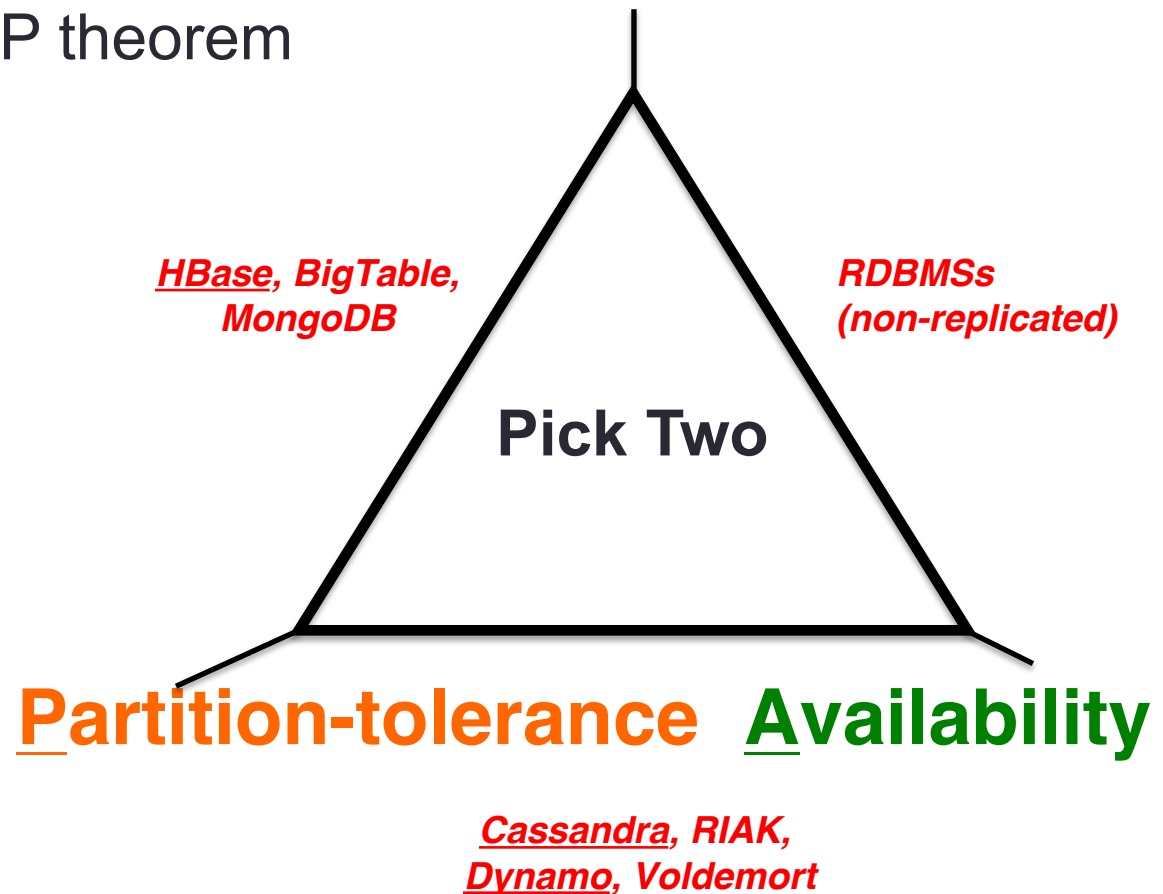
- Strongly Consistent: Read Preference is Master
- Eventually Consistent: Read Preference is Slave (Secondary)
- CAP Theorem: With strong consistency, under partition, MongoDB becomes write-unavailable thereby ensuring consistency

Reading

- <https://docs.mongodb.org/manual/>

Summary

- Traditional Databases (RDBMSs) work with strong consistency, and offer ACID
- Unfortunately, CAP theorem



Acknowledgement

- These slides contain material developed by Indranil Gupta (UIUC), Steve Ko (Buffalo), and Mainak Ghosh (UIUC).